

CASE STUDY:

ANALYSIS AND OPTIMIZATION OF ENERGY CONSUMPTION

- NGOULAYE KEUNGUEU (DIA)
- ISMAEL KONE (DIA)
- DJIBRIL LALEG (DIA)
- LISA NACCACHE (DIA)
- LEINA PRIEUR (DIA)
- KYLIE WU (OCC)
- ILAN ZINI (DIA)

01

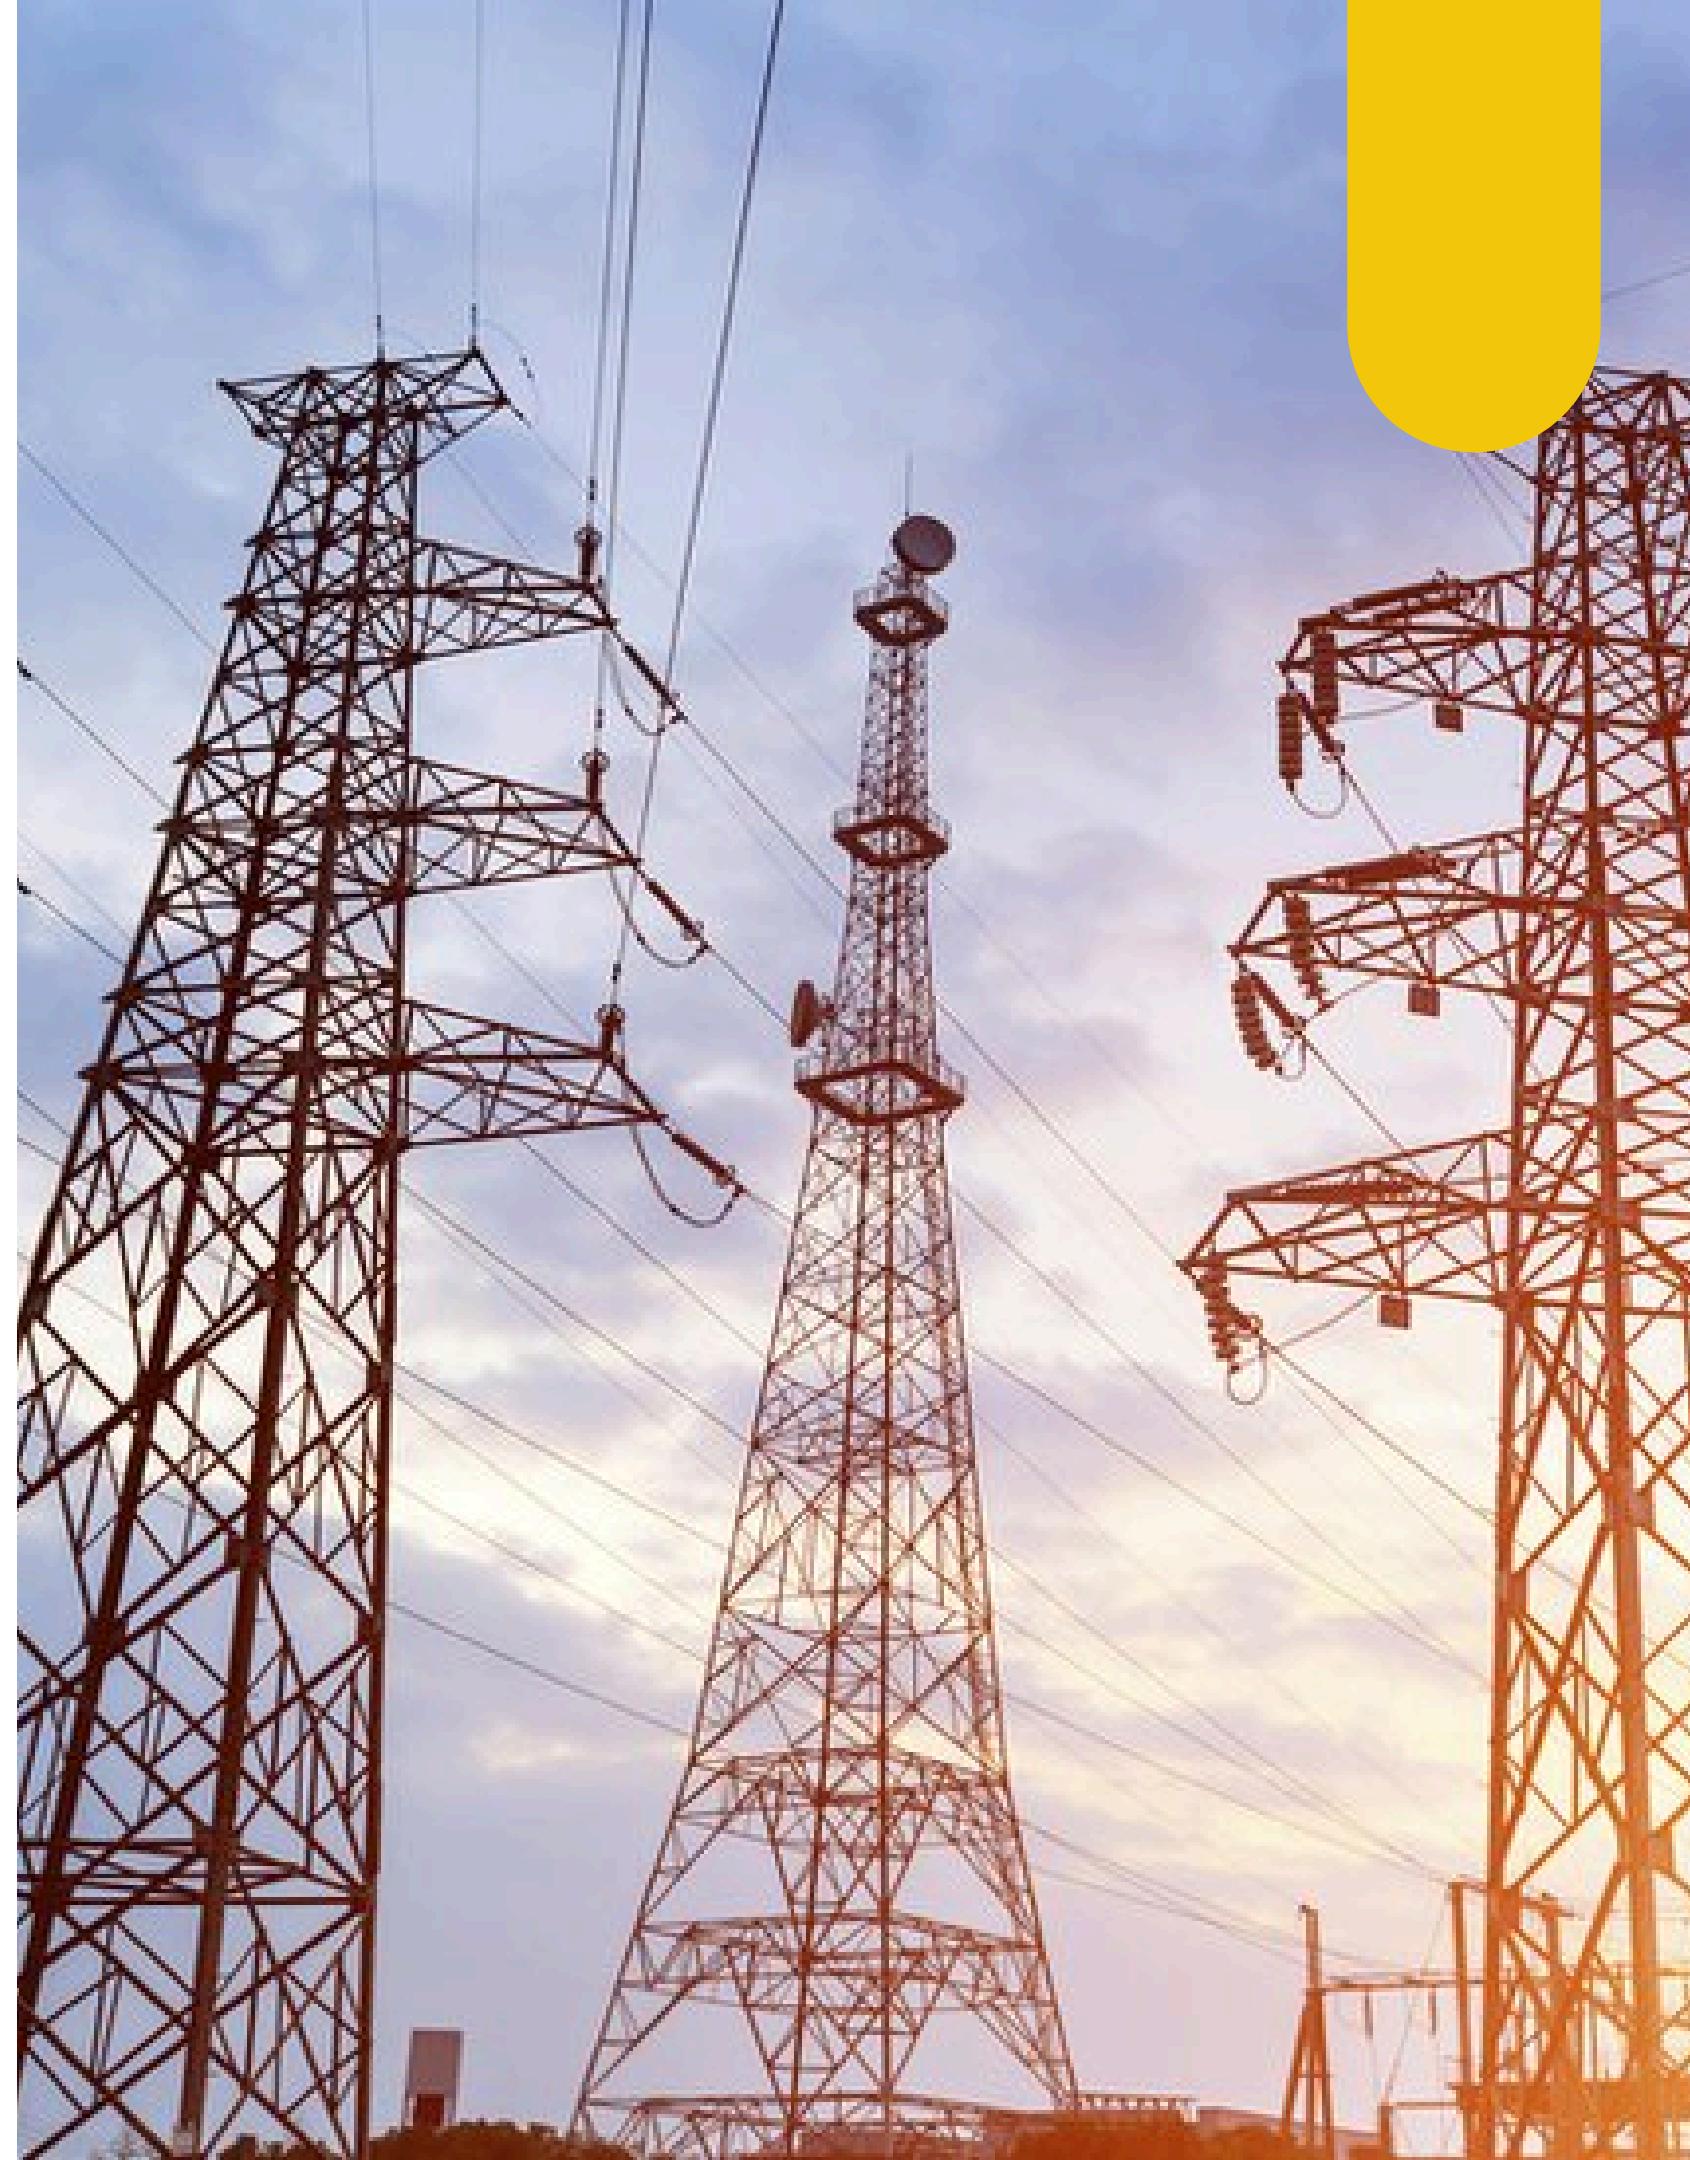


CONTEXT

In the context of the energy transition and global warming, businesses and organizations must minimize their environmental footprint while making efficient use of their resources. Managing energy consumption is a key strategic lever in addressing these global challenges.

Goal: use a dataset to predict the weekly energy consumption of a company.

02



• • • • •
• • • • •
• • • • •
• • • • •

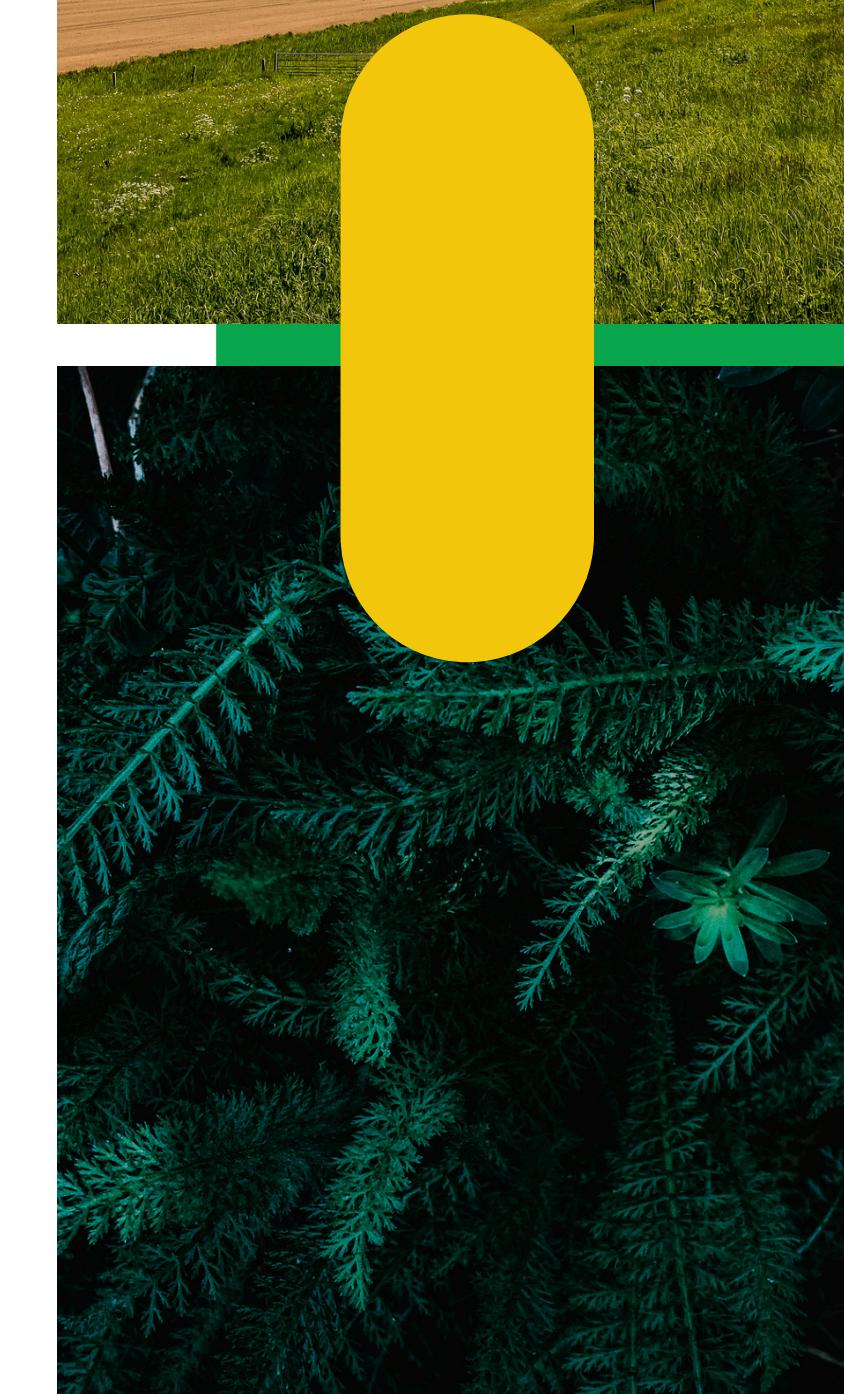
PHASE 1: DATA COLLECTION

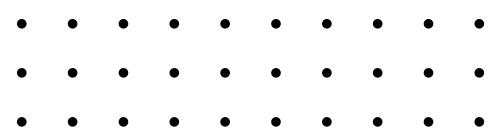
Data collection: our dataset on energy consumption, specifically electricity usage data from 881 companies and local authorities across six French overseas regions : Réunion Island, French Guiana, Martinique, Guadeloupe, Mayotte, and Corsica. The data was collected between 2021 and 2024.

`data.shape`

(42288, 112)

Our dataset contains 42,288 records and 112 columns



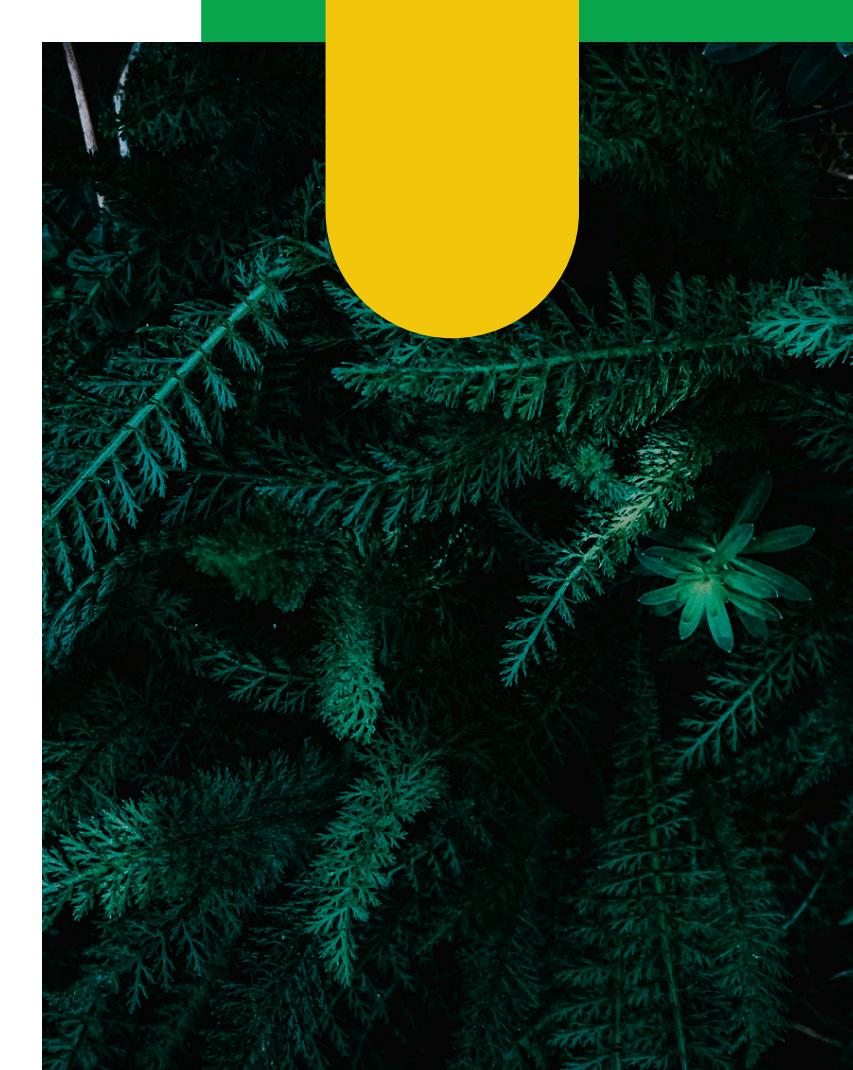


PHASE 1: DATA PREPARATION

Missing values : We decide to delete all columns that have more than 45 % of their values missing. We deleted 90 columns.

We still have a very small pourcentage of missing values for some of our remaining features so we decide to replace the missing values in those features by the mean value of the corresponding feature.

After cleaning, our dataset contains **30,725 rows** and **22 columns**.

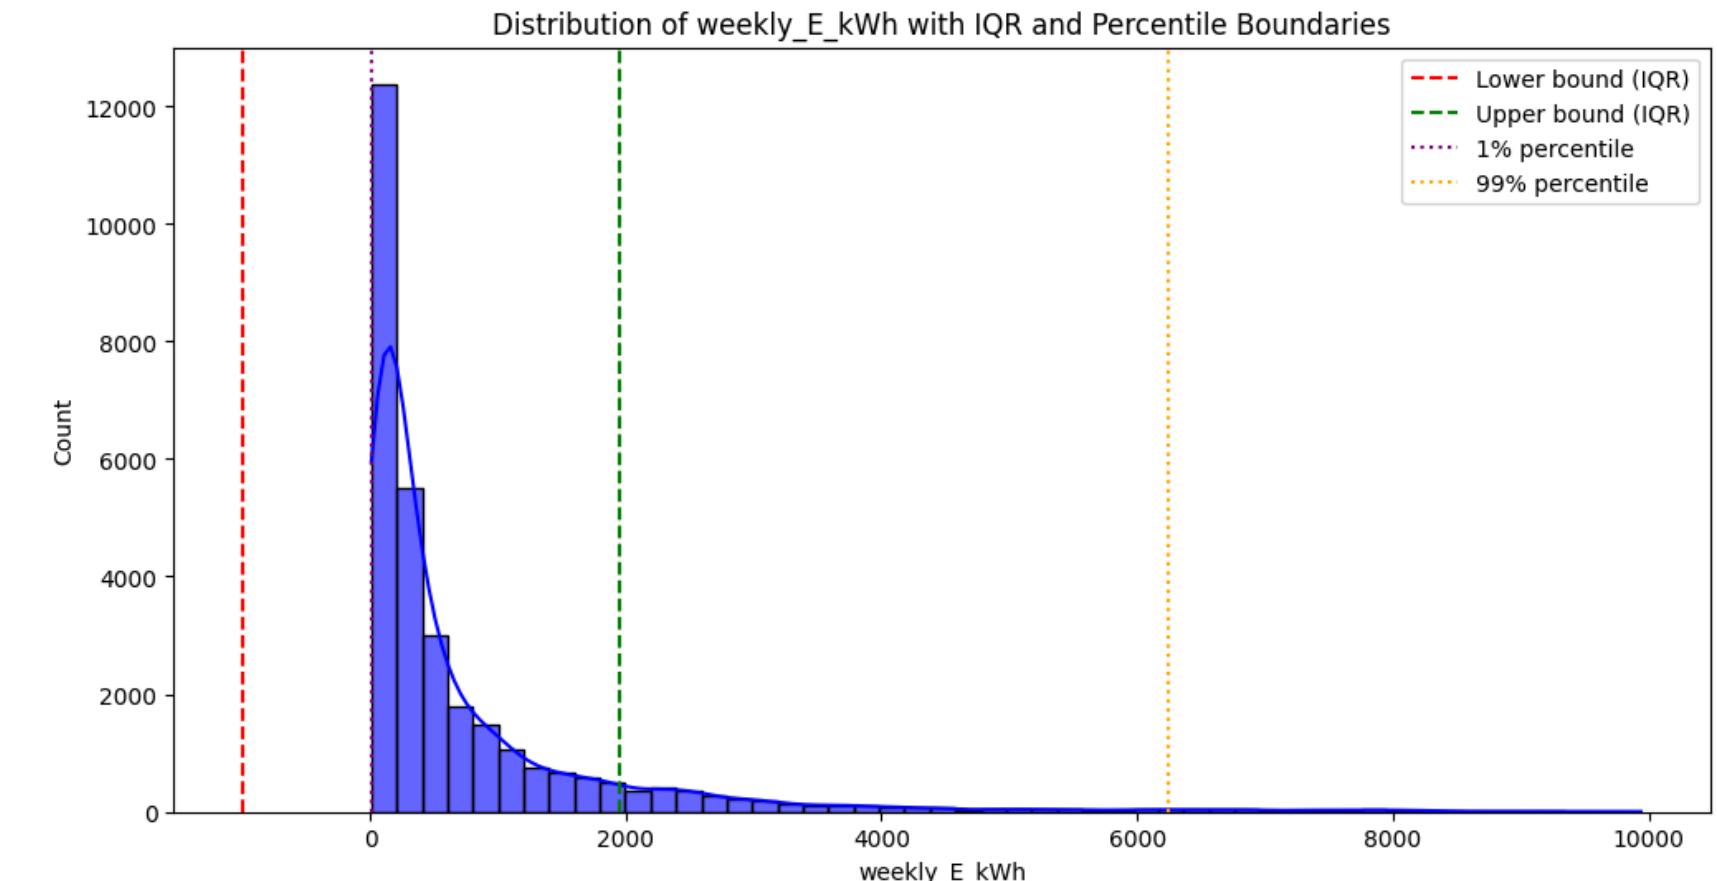
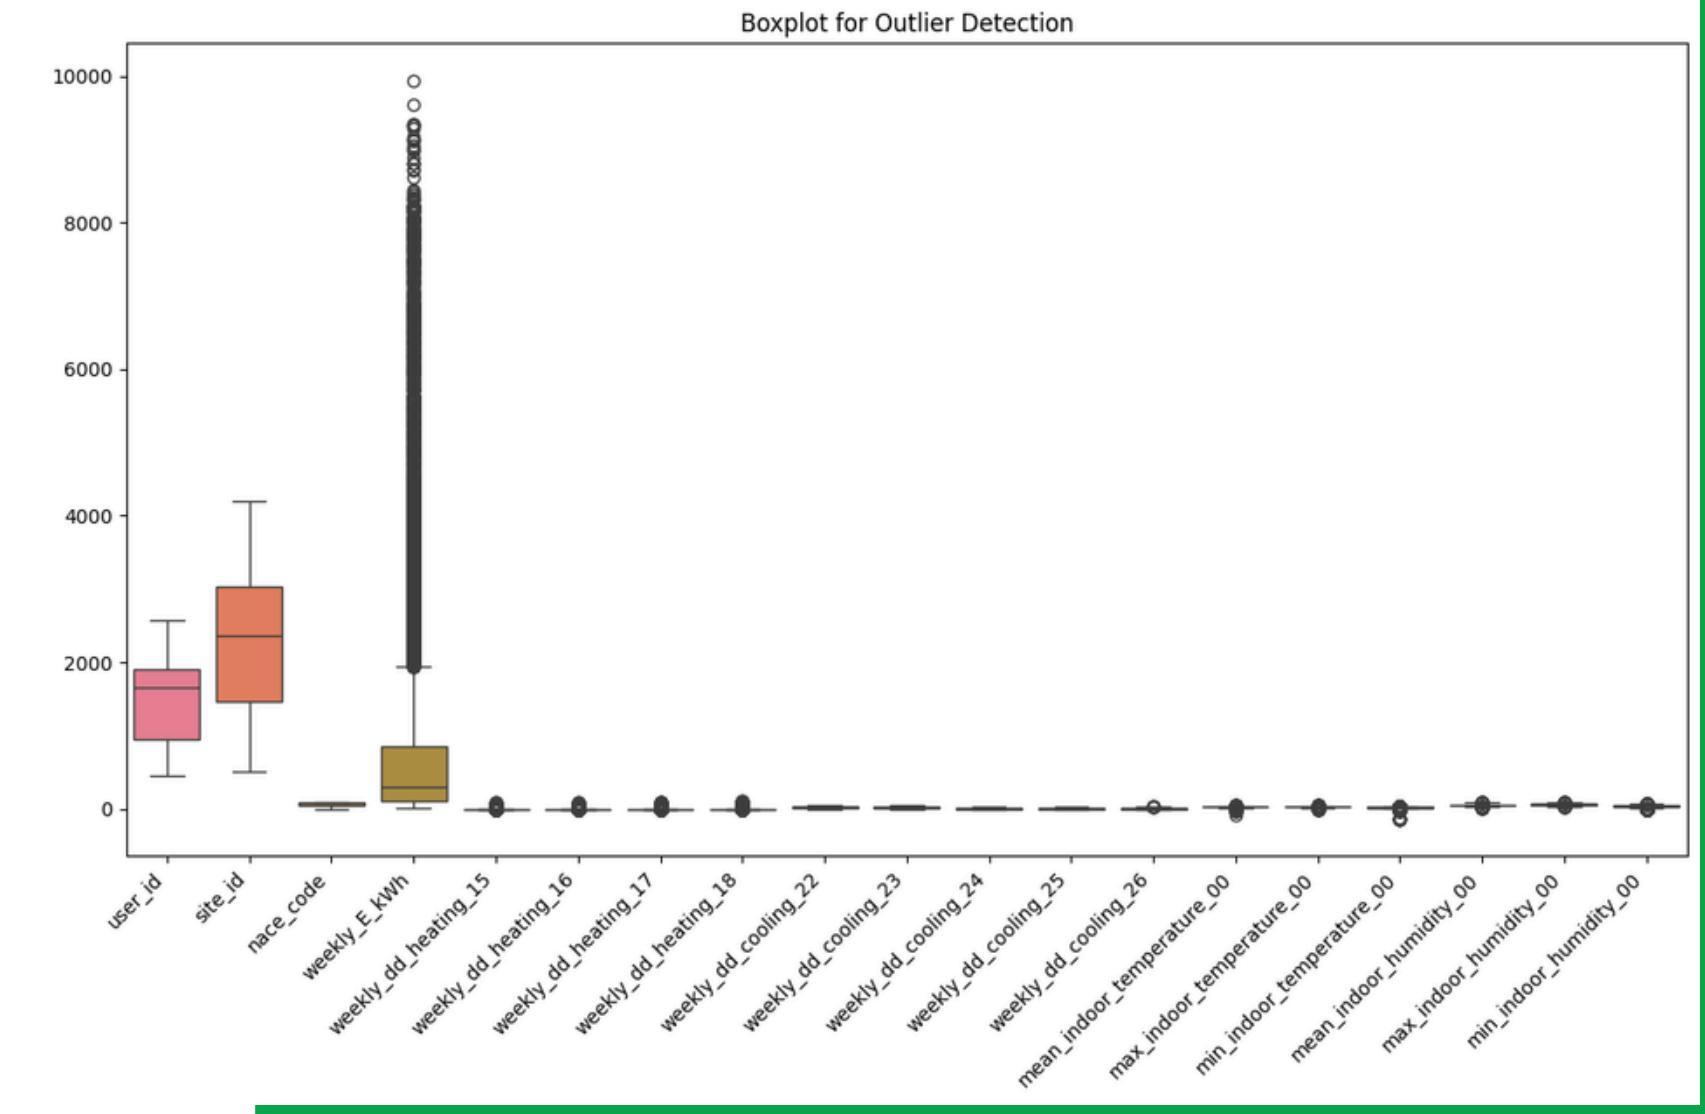


PHASE 1: DATA PREPARATION

Outliers management: We can see that the target column, Weekly Electricity Consumption, contains outliers.

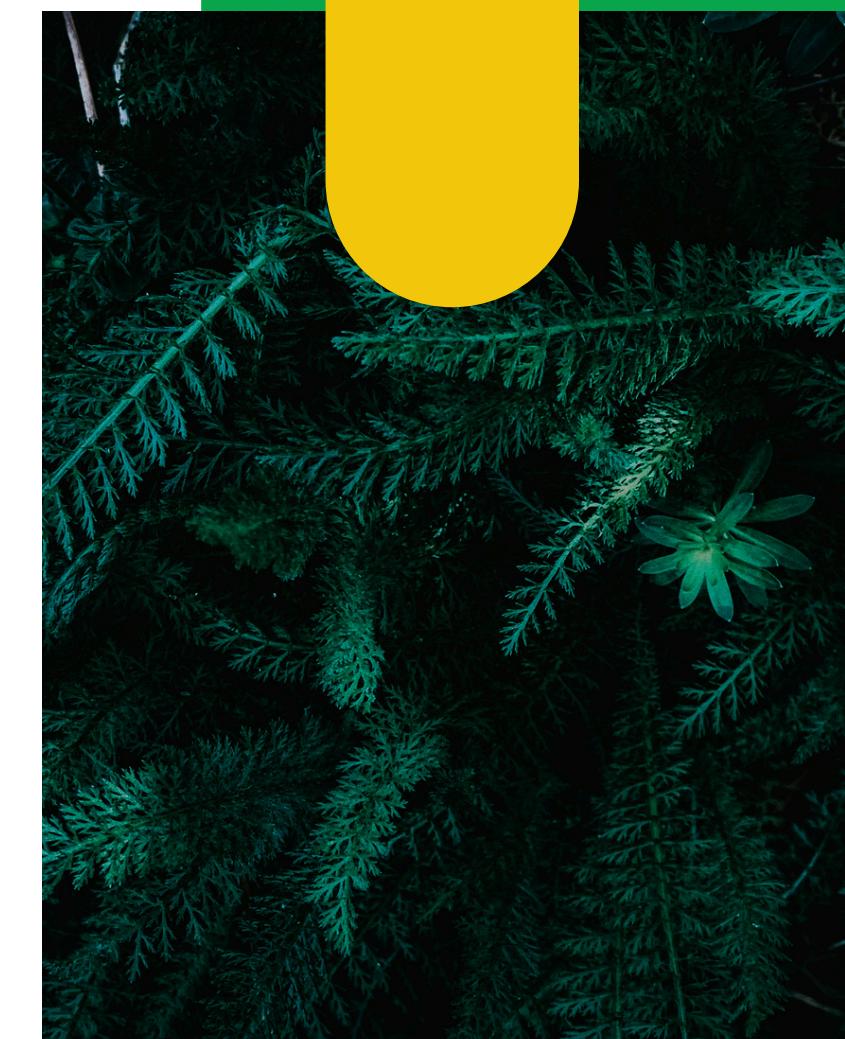
We plotted of the Distribution of Weekly electricity consumption with IQR and Percentile Boundaries.

We decide not to do anything about this since some machine learning models like Decision Trees, Random Forests, and Gradient Boosting (XGBoost, LightGBM, CatBoost) handle outliers well.



PHASE 1: DATA PREPARATION

Feature Engineering / Encoding :



• • • • •
• • • • •
• • • • •
• • • • •

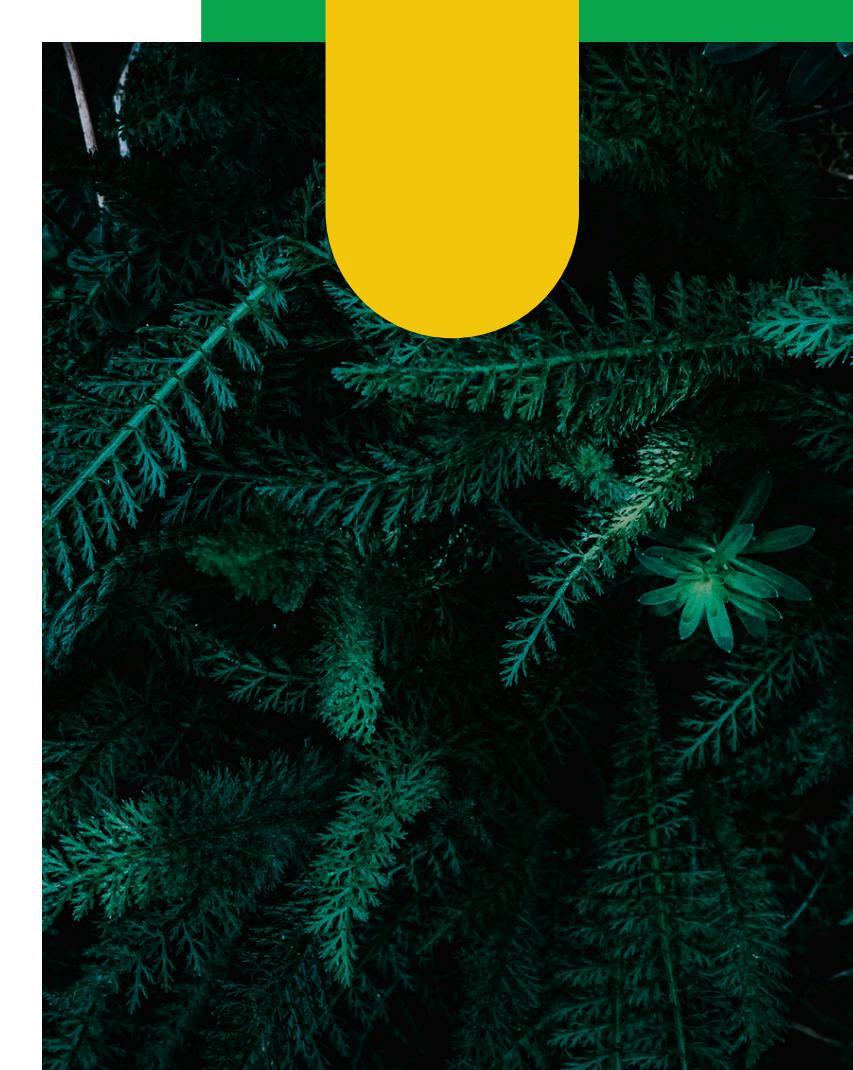
PHASE 1: DATA PREPARATION

Missing values : We decide to delete all columns that have more than 45 % of their values missing.

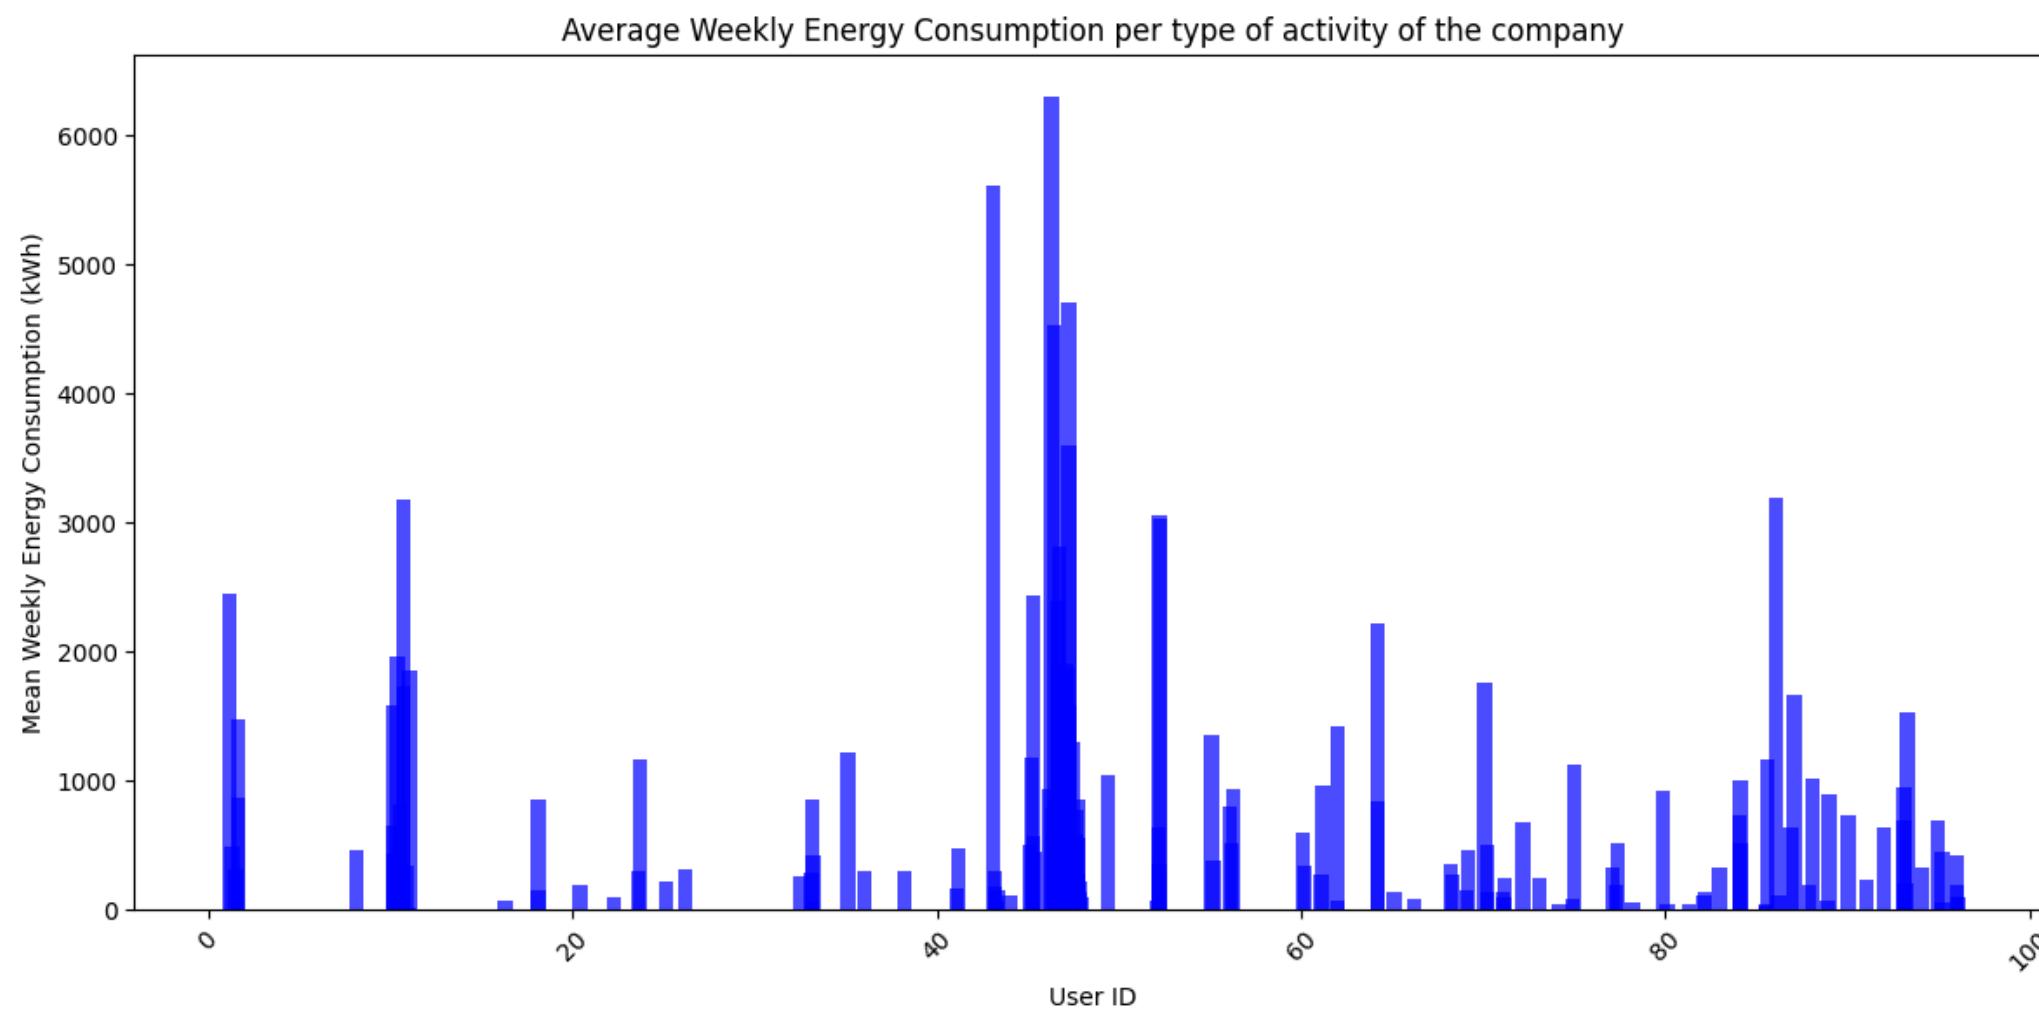
We deleted 90 columns.

We still have a very small pourcentage of missing values for some of our remaining features so we decide to replace the missing values in those features by the mean value of the corresponding feature.

After cleaning, our dataset contains **30,725 rows** and **22 columns**.

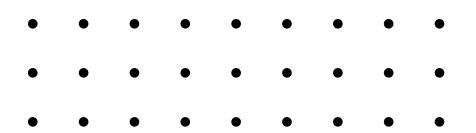


PHASE 2: Analysis and Modeling

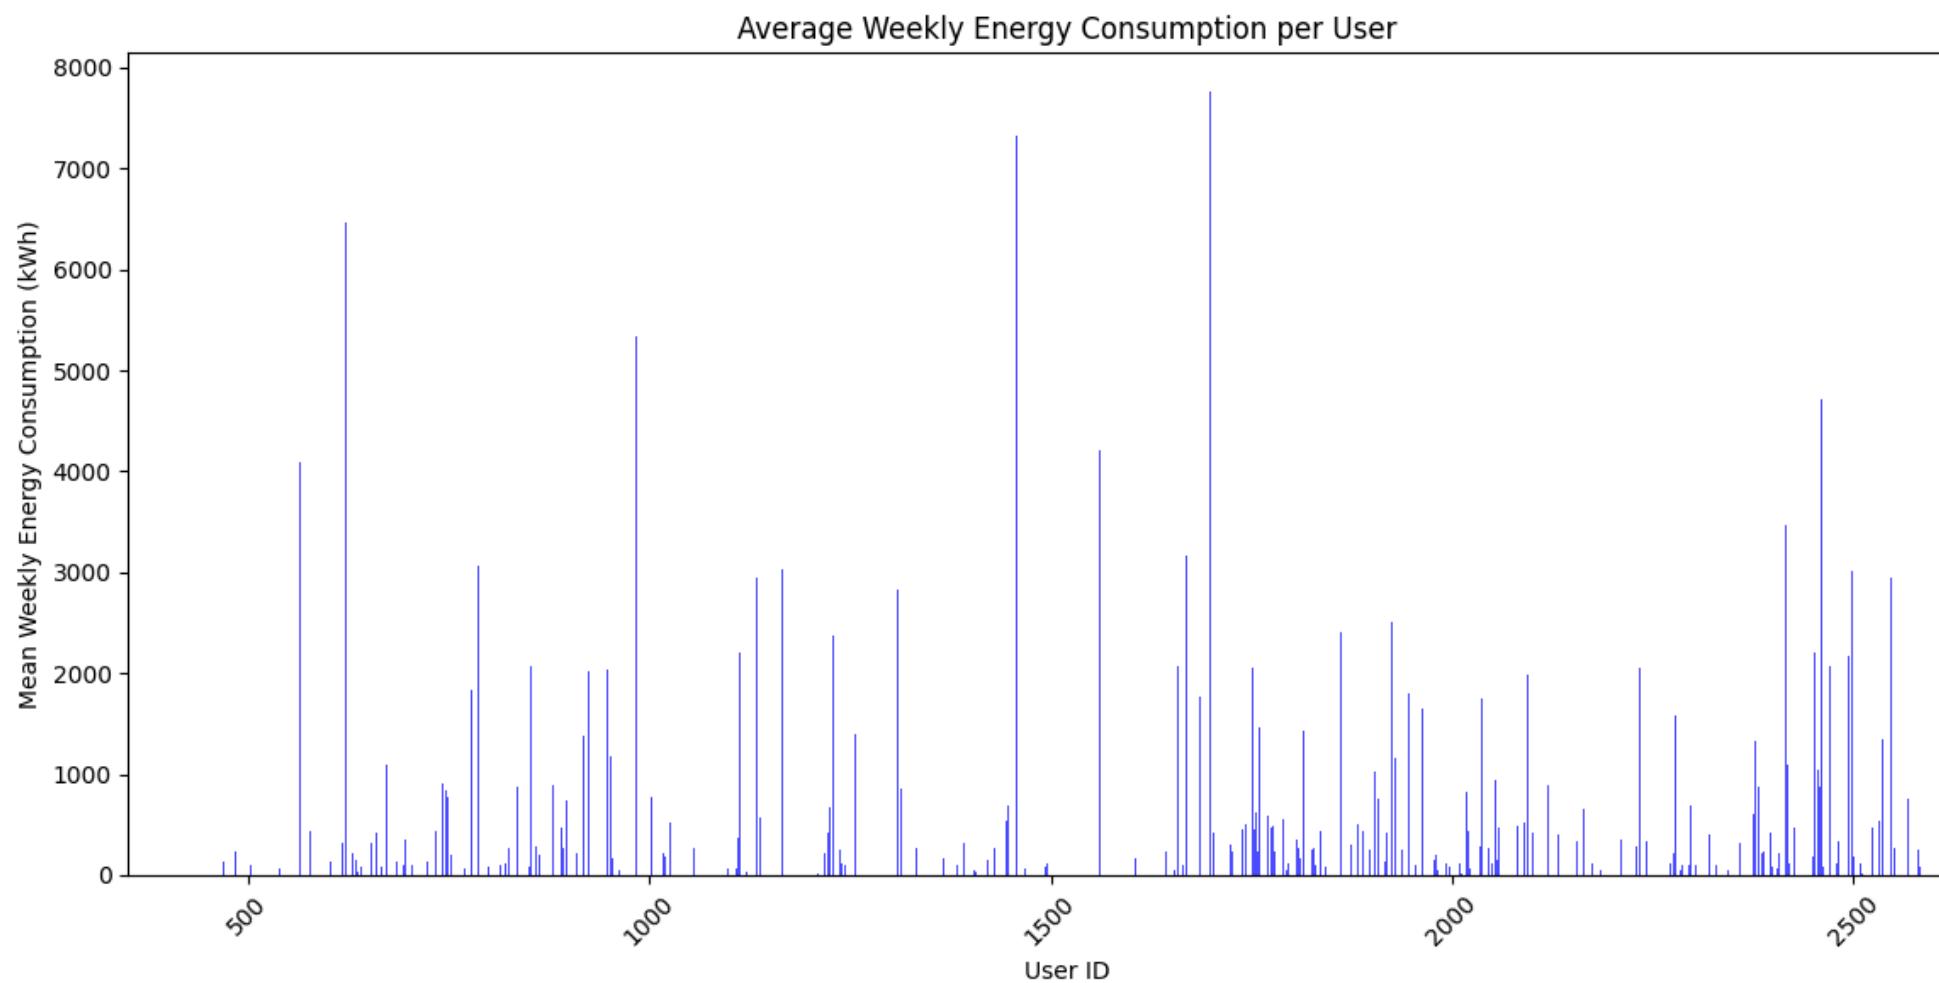


Exploratory data analysis (EDA):
Average weekly energy consumption per type of activity of the company

Observation : Energy consumption varies significantly across company activity types, with a few types showing notably higher average weekly energy use, indicating they are more energy-intensive. Most other activities consume relatively less energy in comparison.



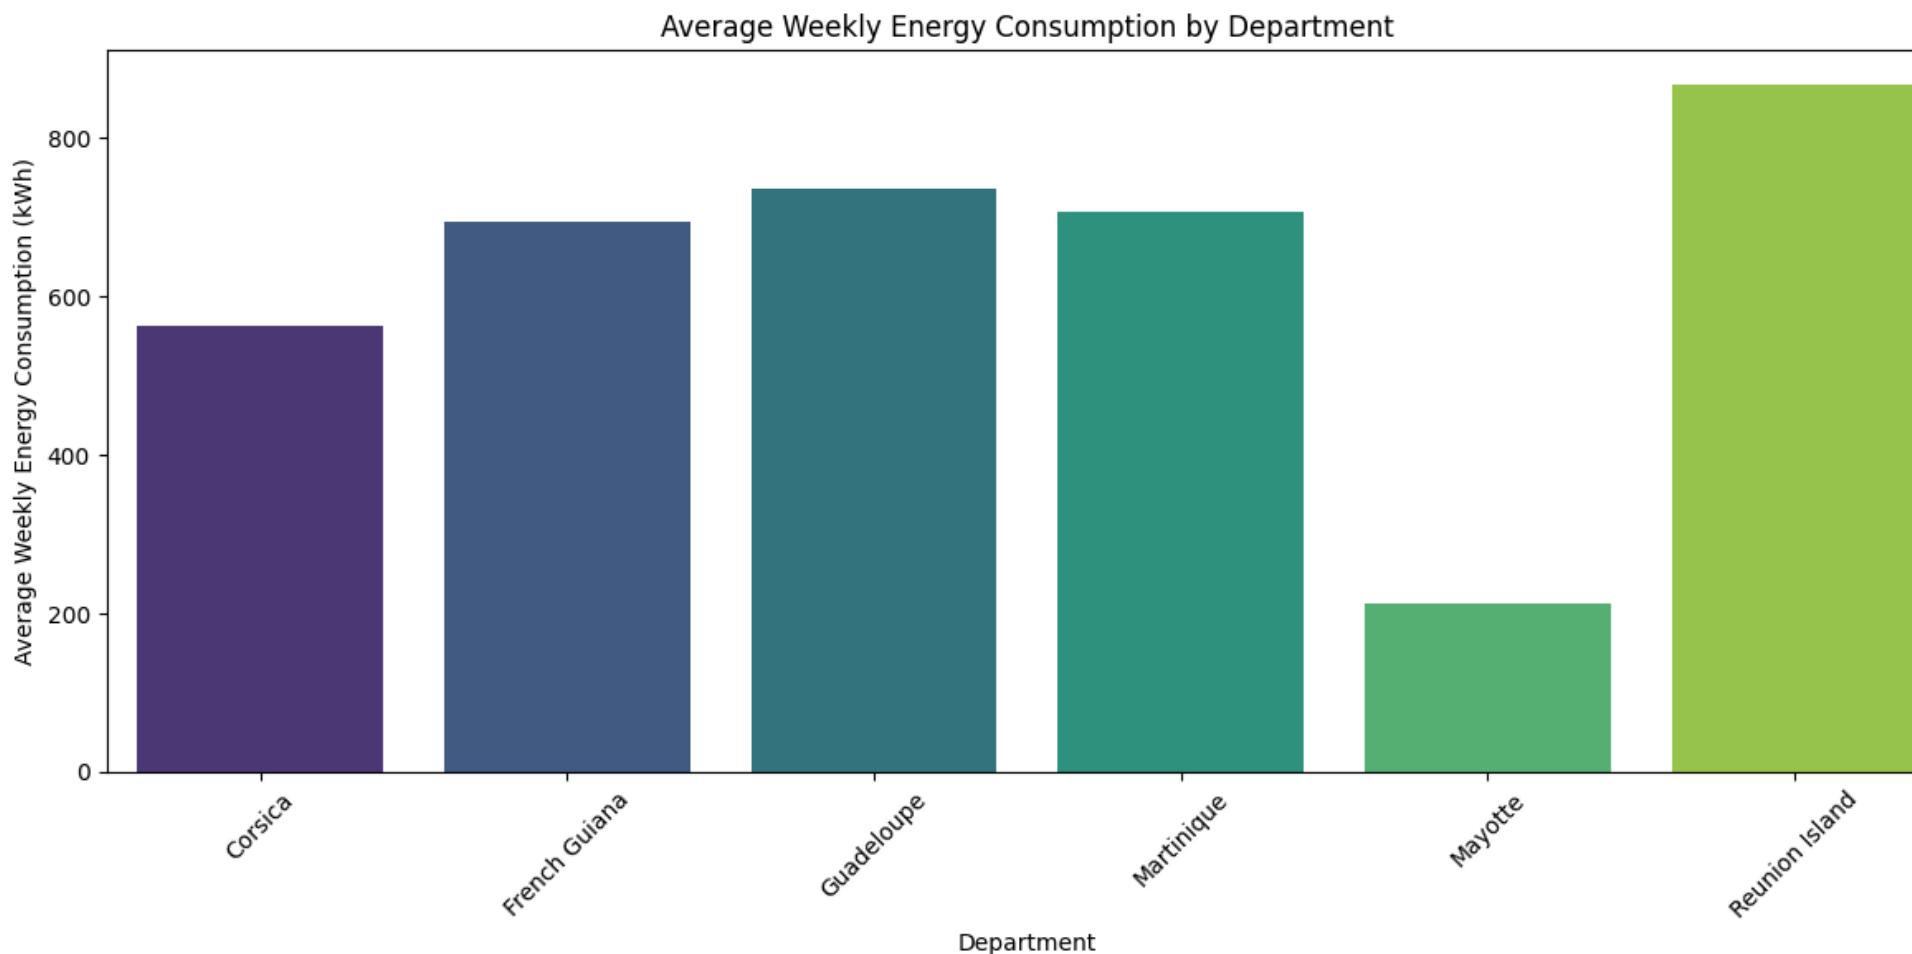
PHASE 2: Analysis and Modeling



Exploratory data analysis (EDA):
Average Weekly Energy Consumption per User

Observation: Most users have relatively low average weekly energy consumption, but there are a few users who consume significantly more energy, indicating high-demand users.

PHASE 2: Analysis and Modeling



Exploratory data analysis (EDA):
Energy consumption by department

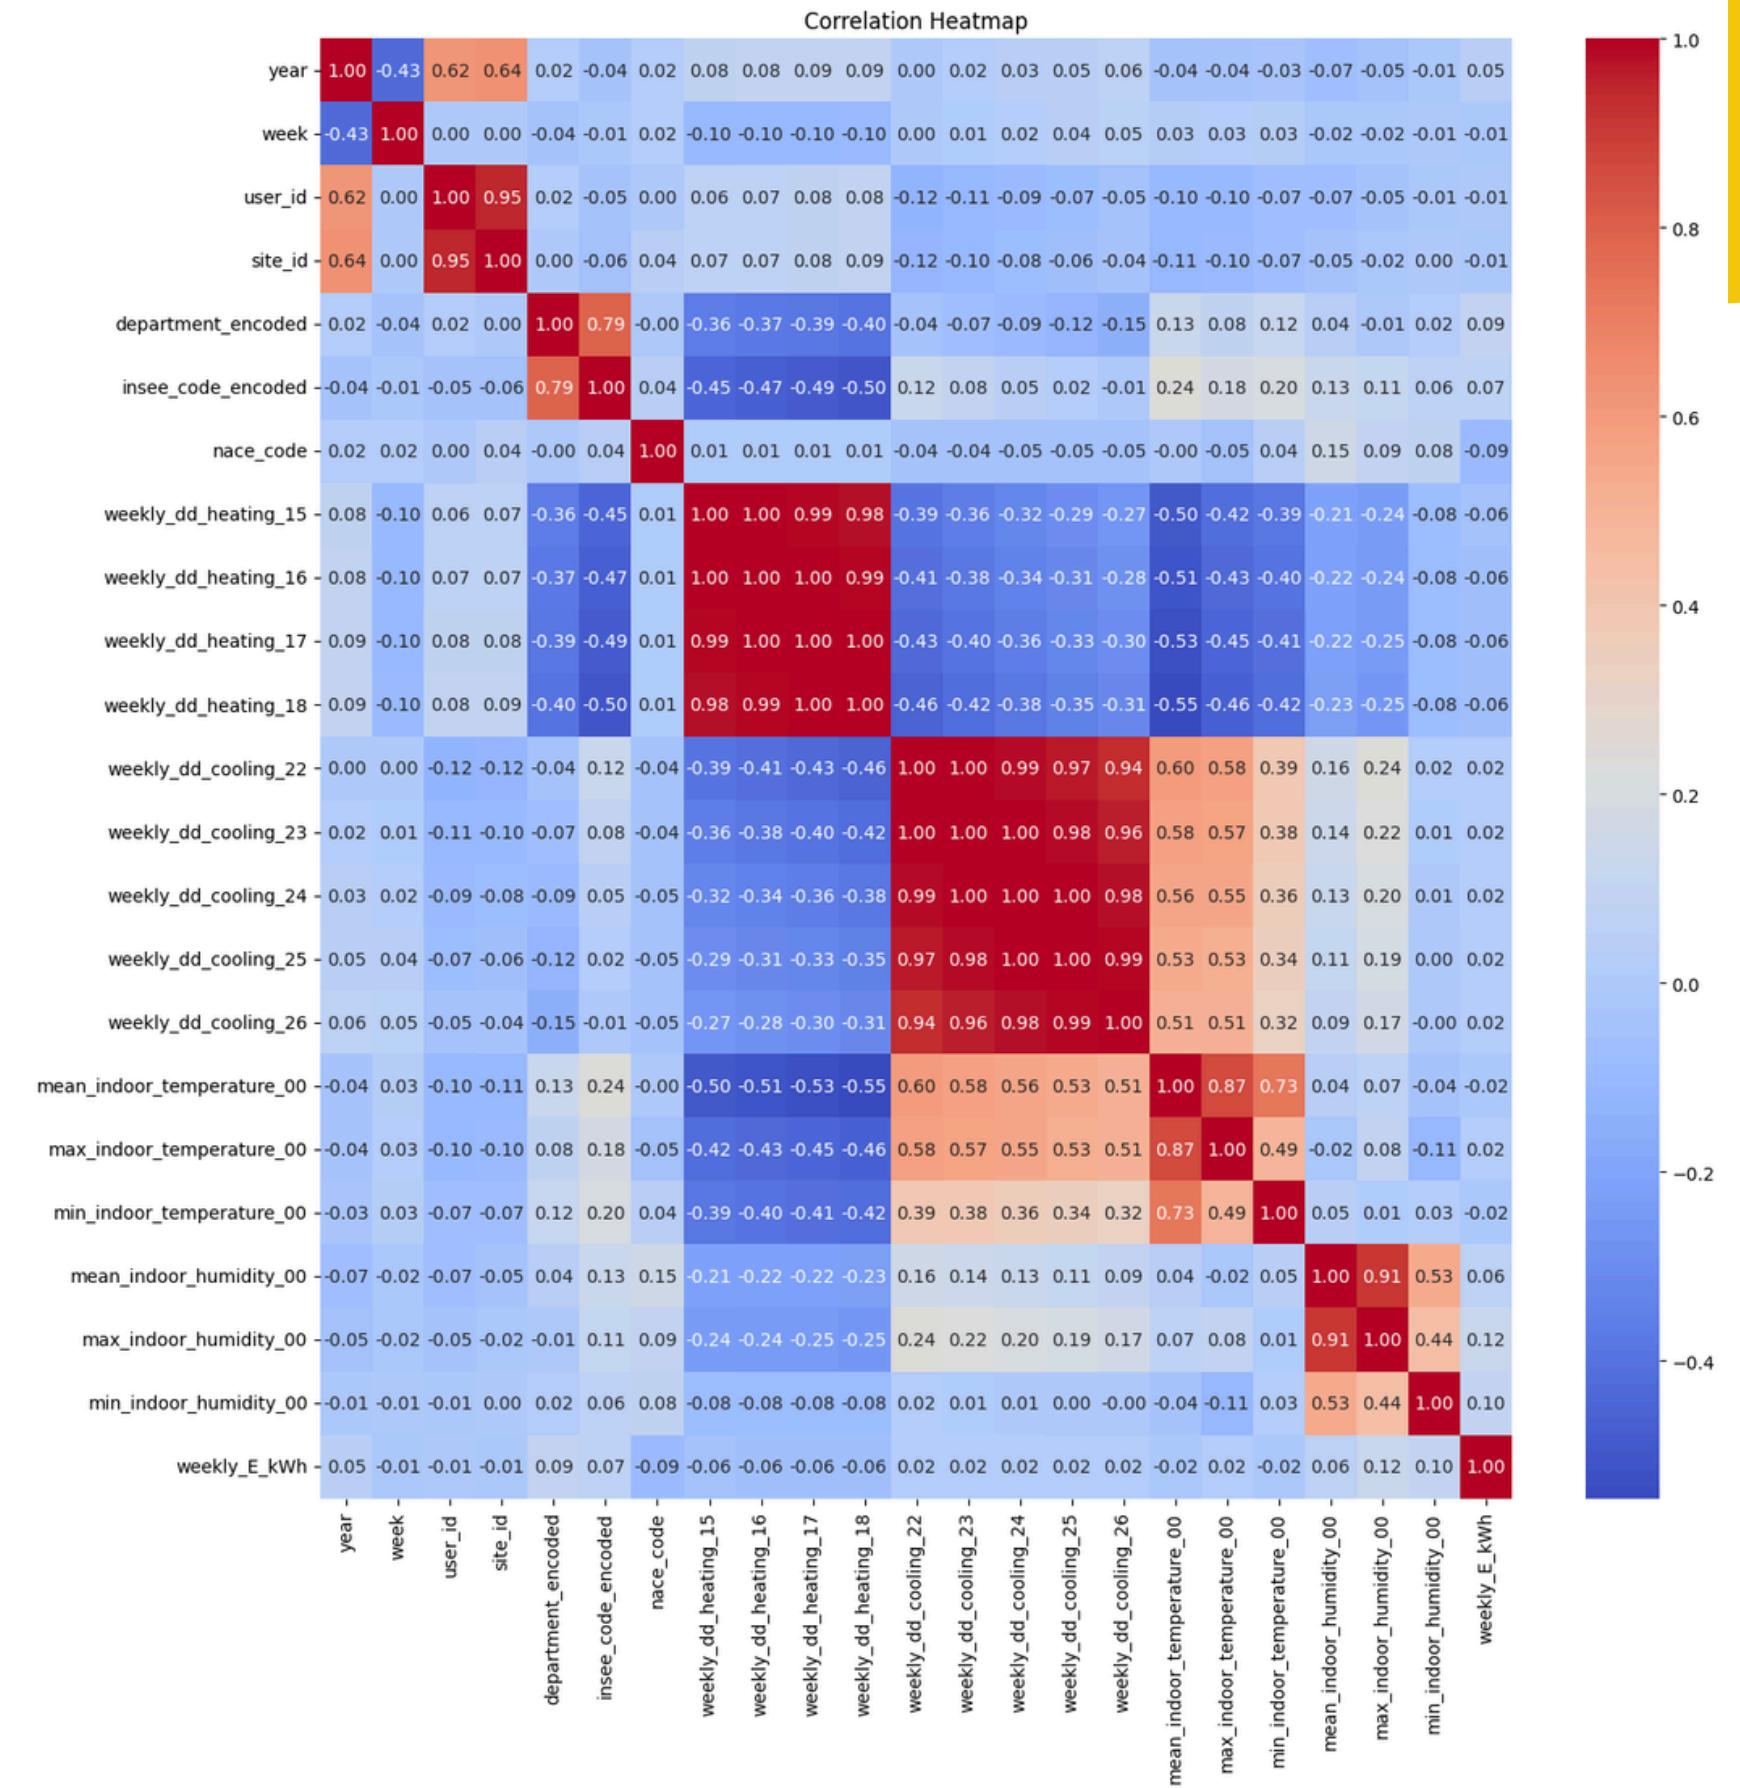
Observation: Clear disparities in average weekly energy consumption among departments. Notably, Réunion Island has the highest average consumption. This suggests that geographic or operational differences may influence energy usage across regions.

• • • • •

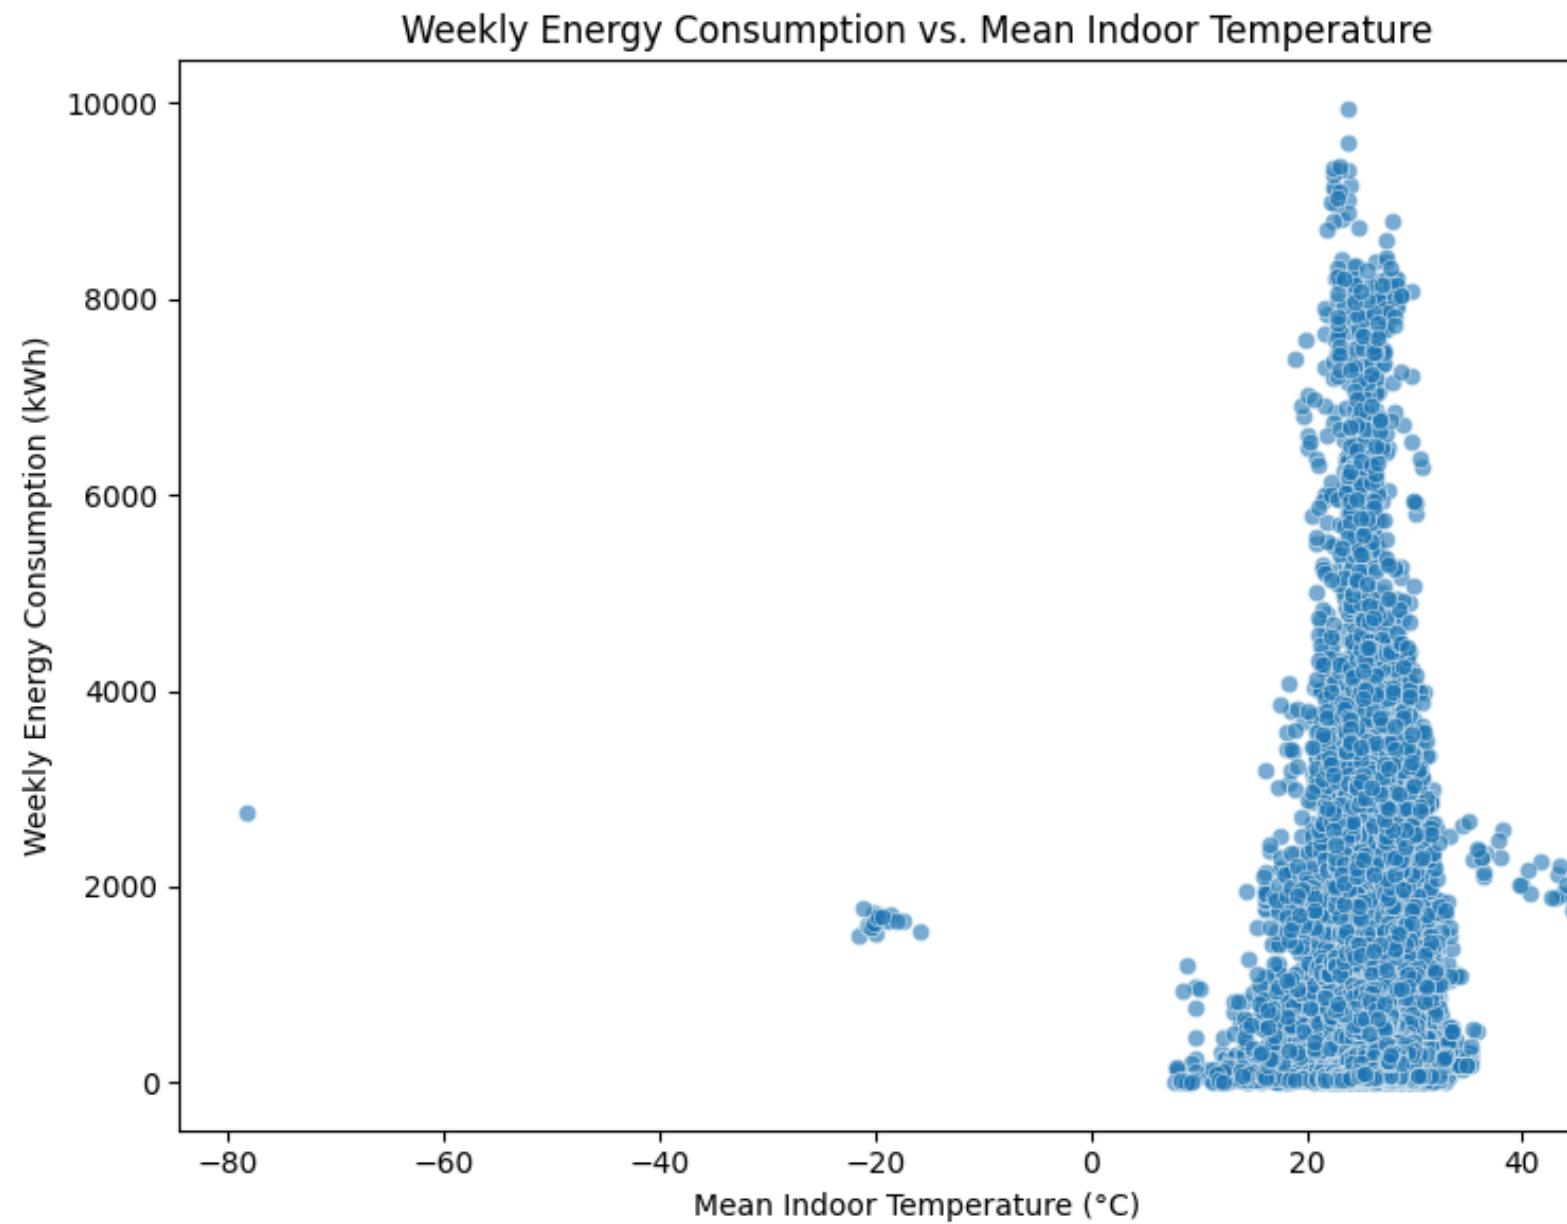
PHASE 2: Analysis and Modeling

Exploratory data analysis (EDA):
Correlation Heatmap

Observation: No significant or noteworthy correlations identified.



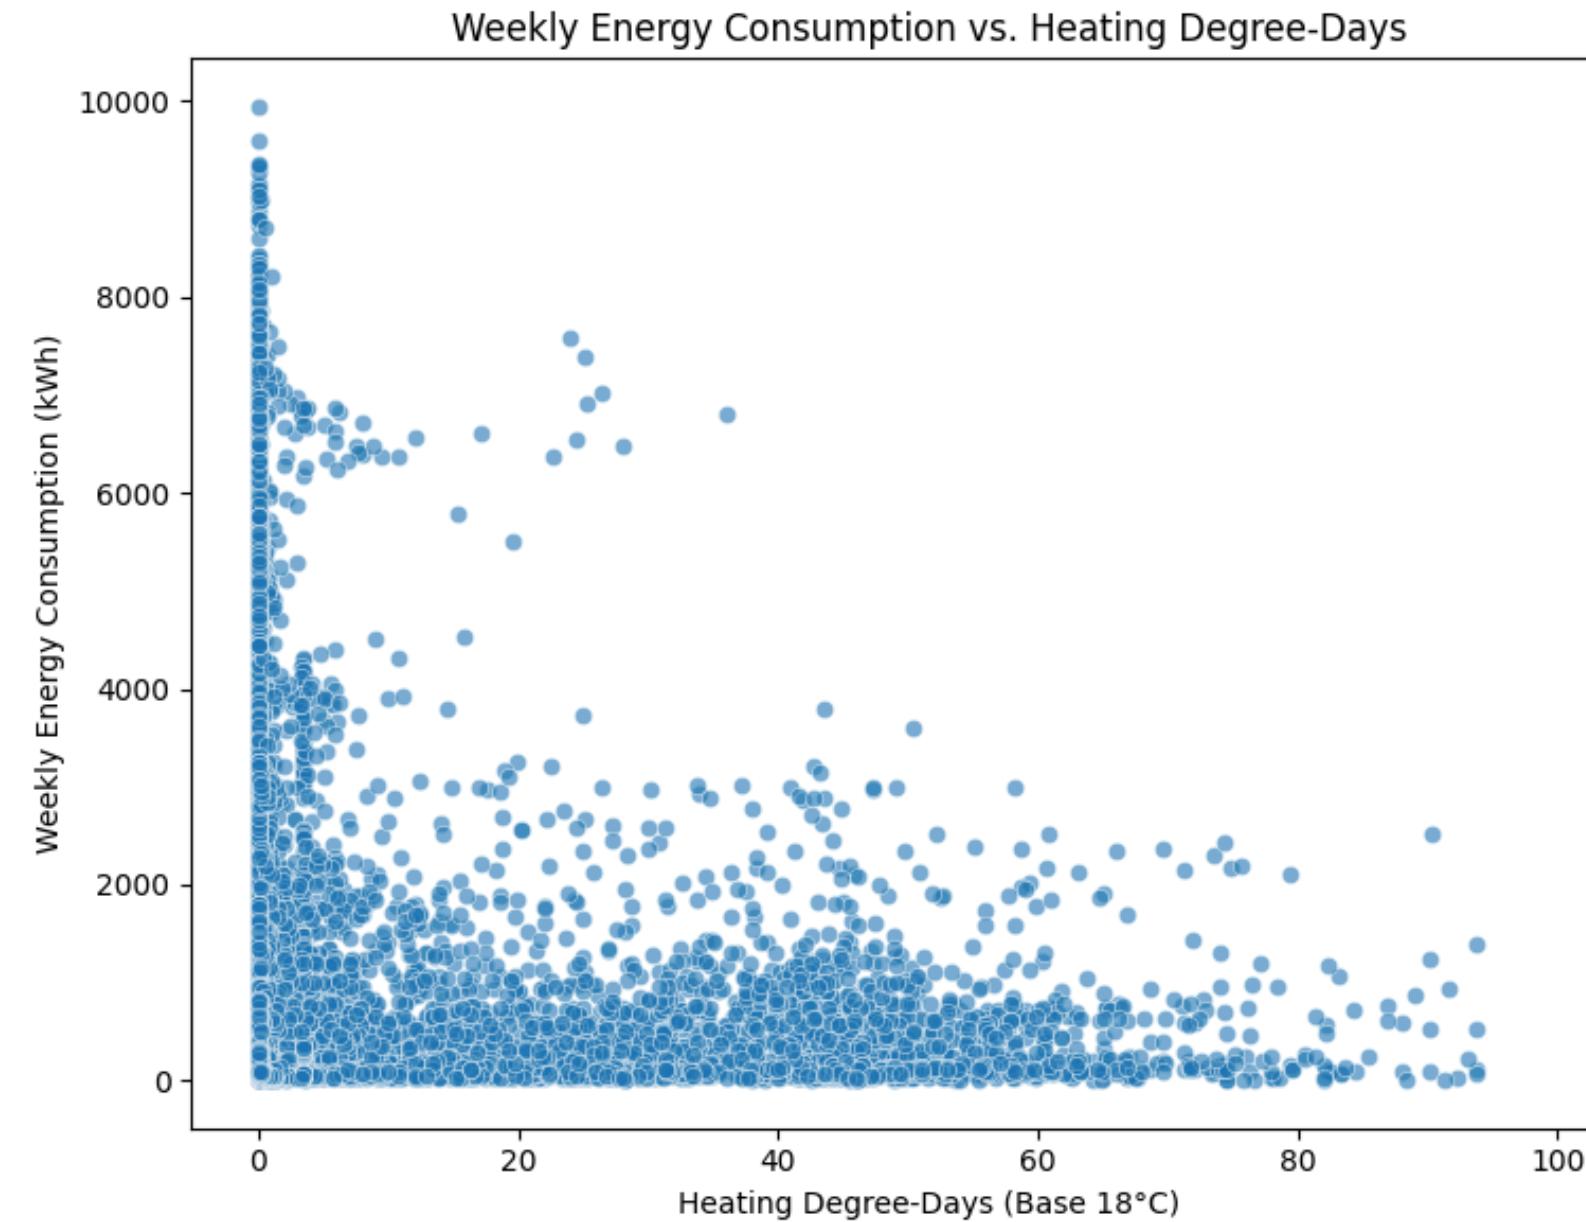
PHASE 2: Analysis and Modeling



Exploratory data analysis (EDA):
Weekly Energy Consumption vs. Mean Indoor Temperature

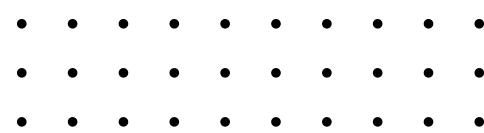
Observation: This scatter plot shows a strong concentration of data between 10°C and 30°C of mean indoor temperature, with most weekly energy consumption values below 4000 kWh.

PHASE 2: Analysis and Modeling



Exploratory data analysis (EDA):
Weekly Energy Consumption vs. Heating Degree-Days

Observation: Shows a weak negative trend, suggesting that energy consumption tends to decrease slightly as heating degree-days increase. A large number of data points cluster around low heating degree-days, indicating that heating is not a dominant factor in overall energy consumption for many users



PHASE 2: Analysis and Modeling

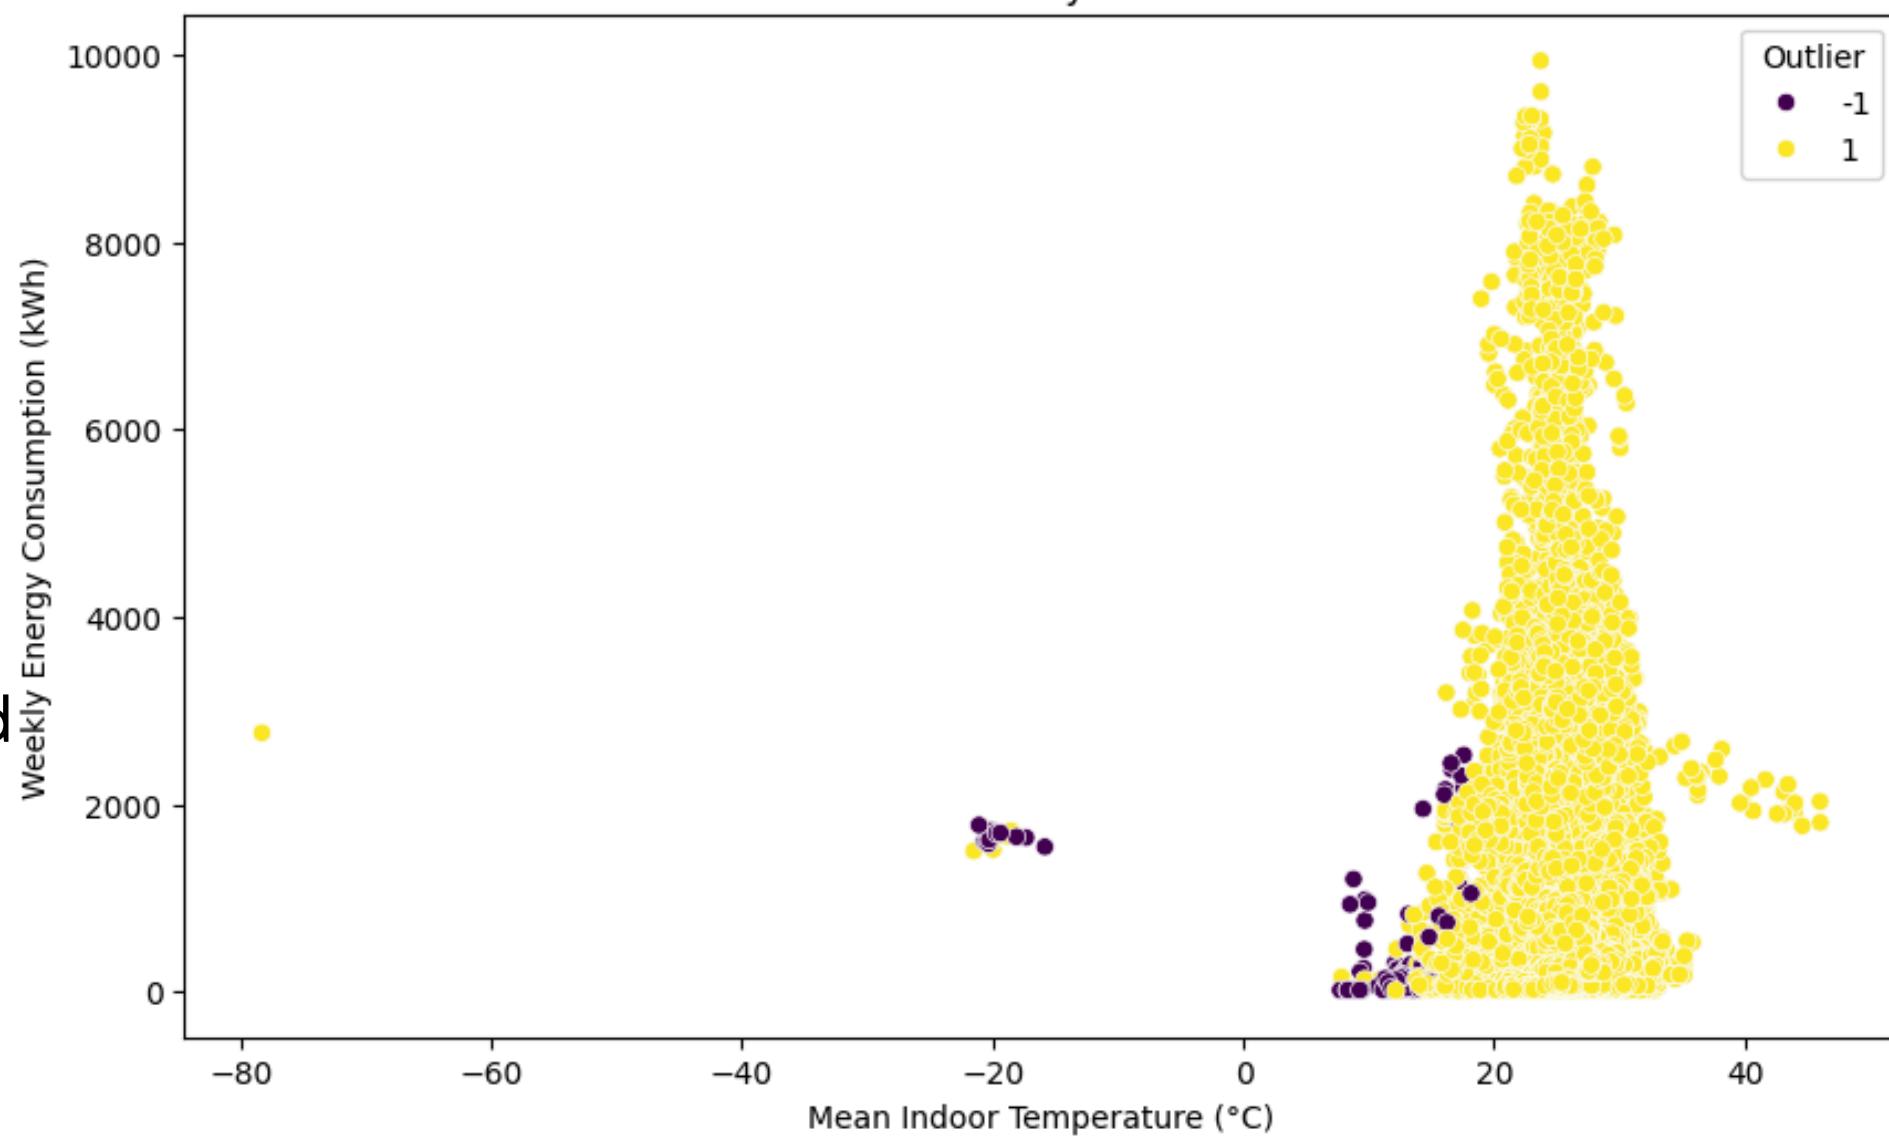
Anomaly Detection for Energy Peaks:

Outliers Detected by Isolation Forest

Observation: Most data points (in yellow) are considered normal, while the algorithm identified a cluster of anomalous points (in purple), mainly located at the extremes of the mean indoor temperature axis, especially below 0°C and around 20°C with unusually low or high energy consumption. These outliers likely represent abnormal temperature readings or unusual consumption behavior, and may indicate data entry errors or special operational conditions worth further investigation



Outliers Detected by Isolation Forest

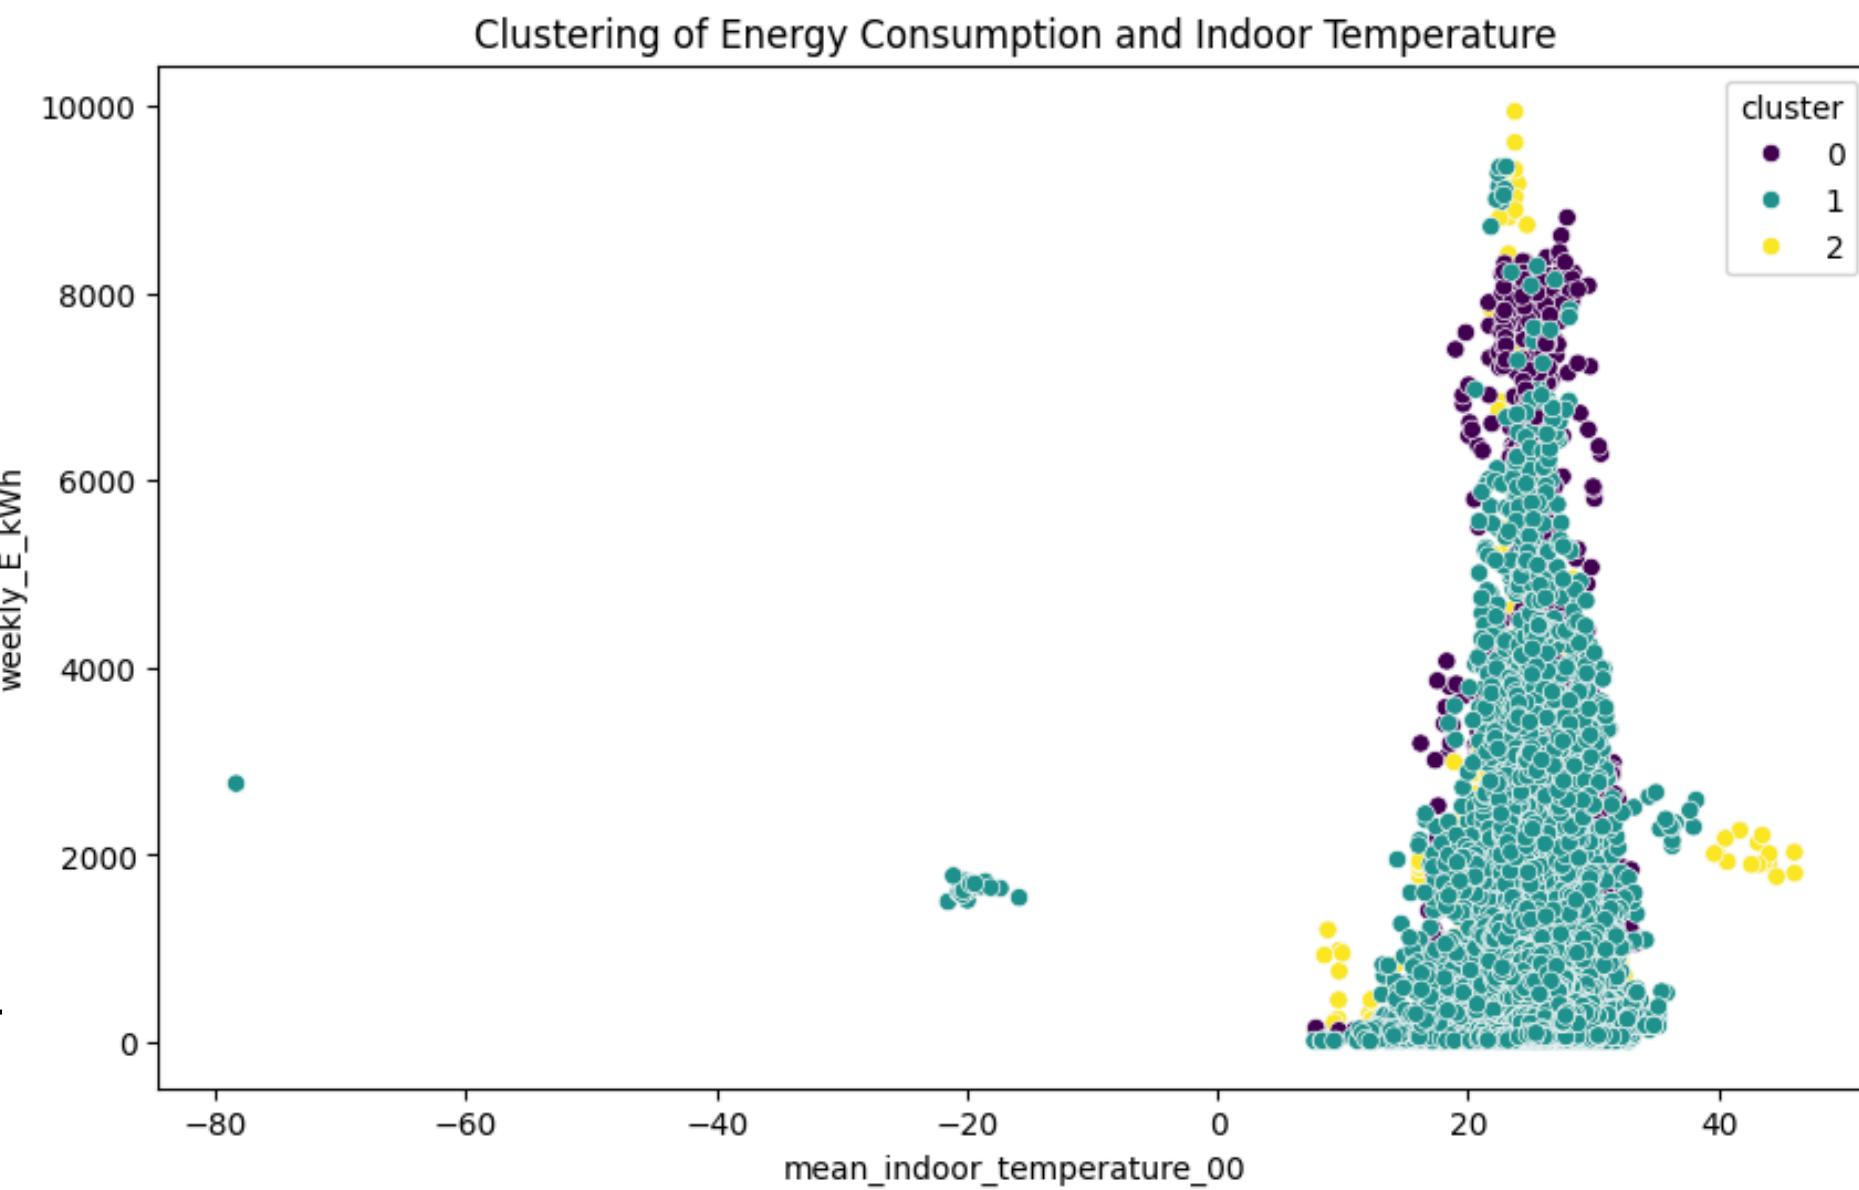


PHASE 2: Analysis and Modeling

Anomaly Detection for Energy Peaks:

Clustering of Energy Consumption and Indoor Temperature using K-Means Clustering

Observation: three distinct groups based on energy consumption and indoor temperature. Most data falls into a central cluster, while smaller clusters highlight high consumption patterns and anomalous temperature values, potentially indicating unusual behaviors.



PHASE 2: Analysis and Modeling

Predictive model for energy consumption:

Model	MAE	RMS	R2	REASON
1 Random Forest Regression	100.9	56534.3	0.96	Robustness, ability to handle non-linear relationships, and resistance to overfitting
XGBoost Regressor	334.4	546.9	0.77	High performance and efficiency, particularly with structured tabular data
LightGBM Regressor	188.6	348.2	0.91	Optimization for speed and scalability, making it well-suited for large datasets with numerous features
Bayesian Ridge Regression	690.2	1112.2	0.06	Ability to capture uncertainty in predictions and its built-in regularization
Neural Networks	598.7	1246.5	-0.17	Modeling complex, non-linear patterns

PHASE 2: Analysis and Modeling

Anomaly Detection for Energy Peaks:

Model	Accuracy	F1-score(0)	F1-score(1)	REASON
Random Forest	0.97	0.98	0.93	Strong performance, robustness, and ability to model complex, non-linear relationships without overfitting
XGB Classifier	0.98	0.99	0.97	Speed, scalability, and excellent results on structured data
SVM Classifier	0.75	0.86	0.0	Effectiveness in finding optimal decision boundaries, particularly in high-dimensional spaces.
KNeighbors Classifier	0.98	0.98	0.95	A simple and intuitive model that works well when similar instances share the same label
Naive Bayes Classifier	0.63	0.73	0.39	Useful for handling categorical features and assuming feature independence.

1

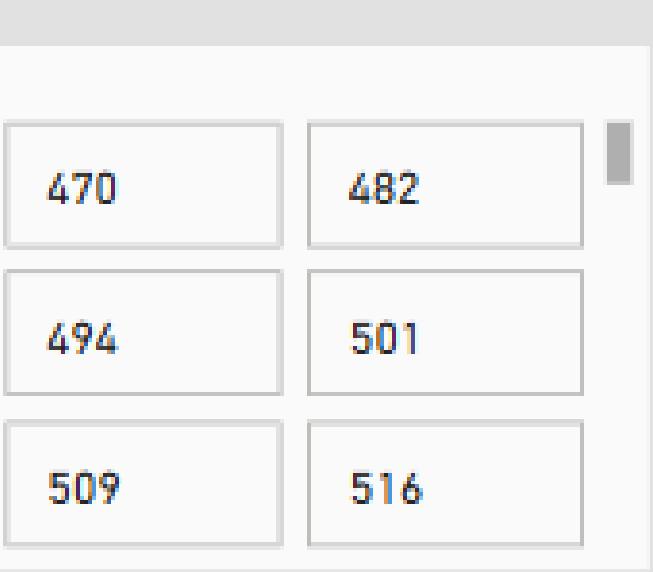
PHASE 2: Analysis and Modeling

Cross-validation :

	Cross-validation	Improvement
Random Forest Regression	Mean Cross-Validation RMSE: 1287.6	
XGB Classifier	Mean Cross-Validation Accuracy: 0.97	Class imbalance : SMOTE + XGBoost Classifier -> 0.99 precision, recall, and F1-score on both classes

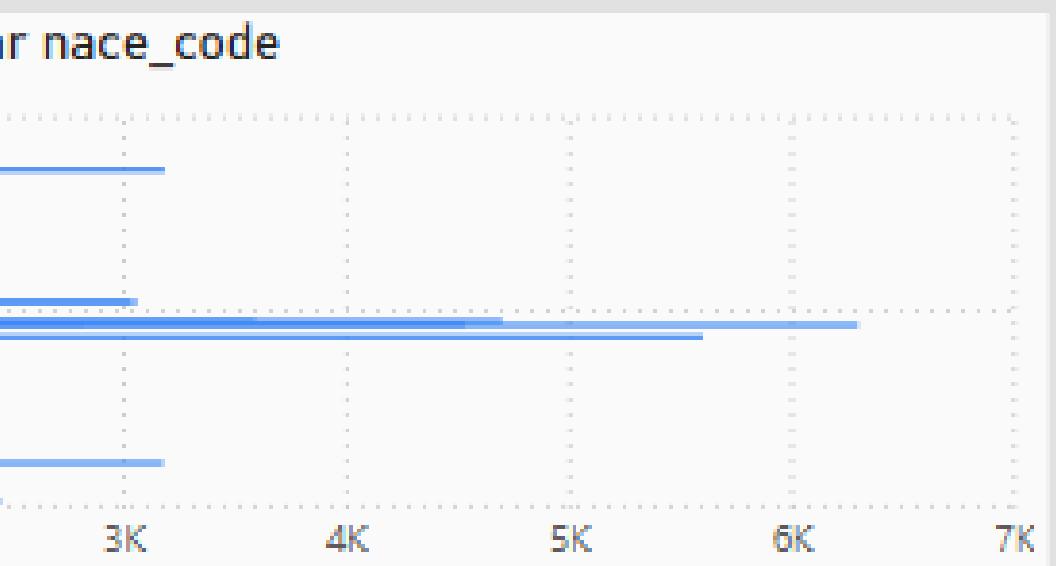
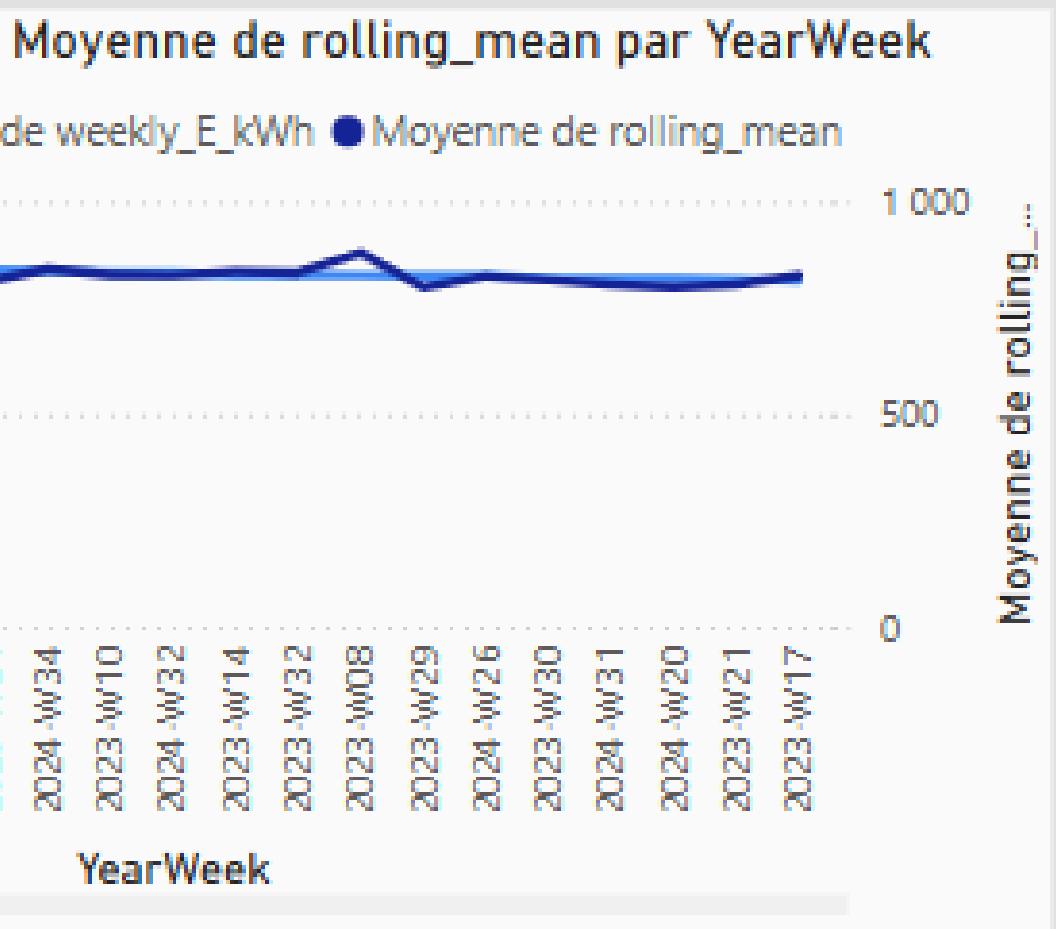


• • • • •
• • • • •
• • • • •
• • • • •

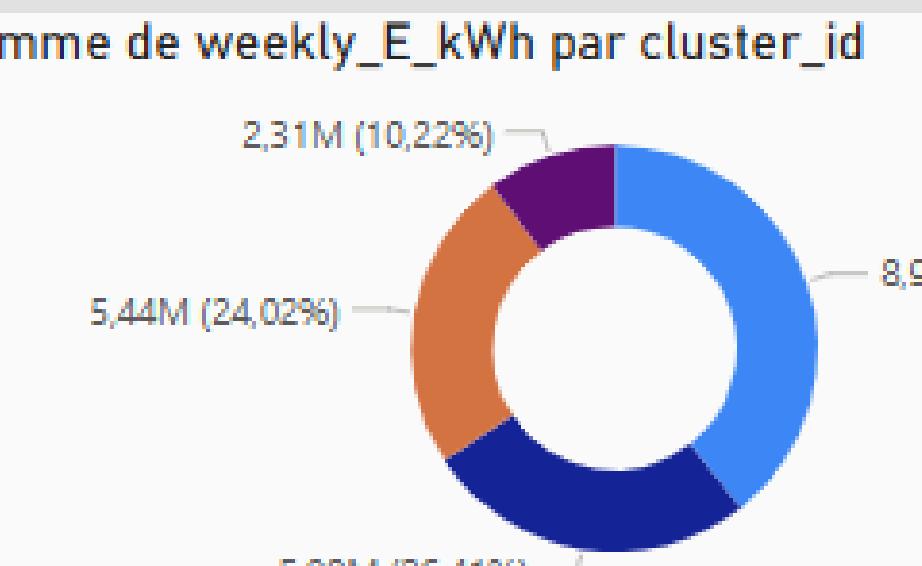


PHASE 3: CREATING A DASHBOARD

To visualize and analyze energy consumption trends, peaks, and patterns, we used Power BI to create an interactive dashboard based on the cleaned and processed dataset.



Dynamic peak table (peak)			
user_id	YearWeek	weekly_E_kWh	nace_code
501	2023-W06	6 950	FR
501	2023-W07	7 350	ES
501	2023-W08	7 080	DE
501	2023-W09	6 576	IT
501	2023-W10	4 295,98	PT
501	2023-W20	4 162,00	NL



PHASE 3: CREATING A DASHBOARD

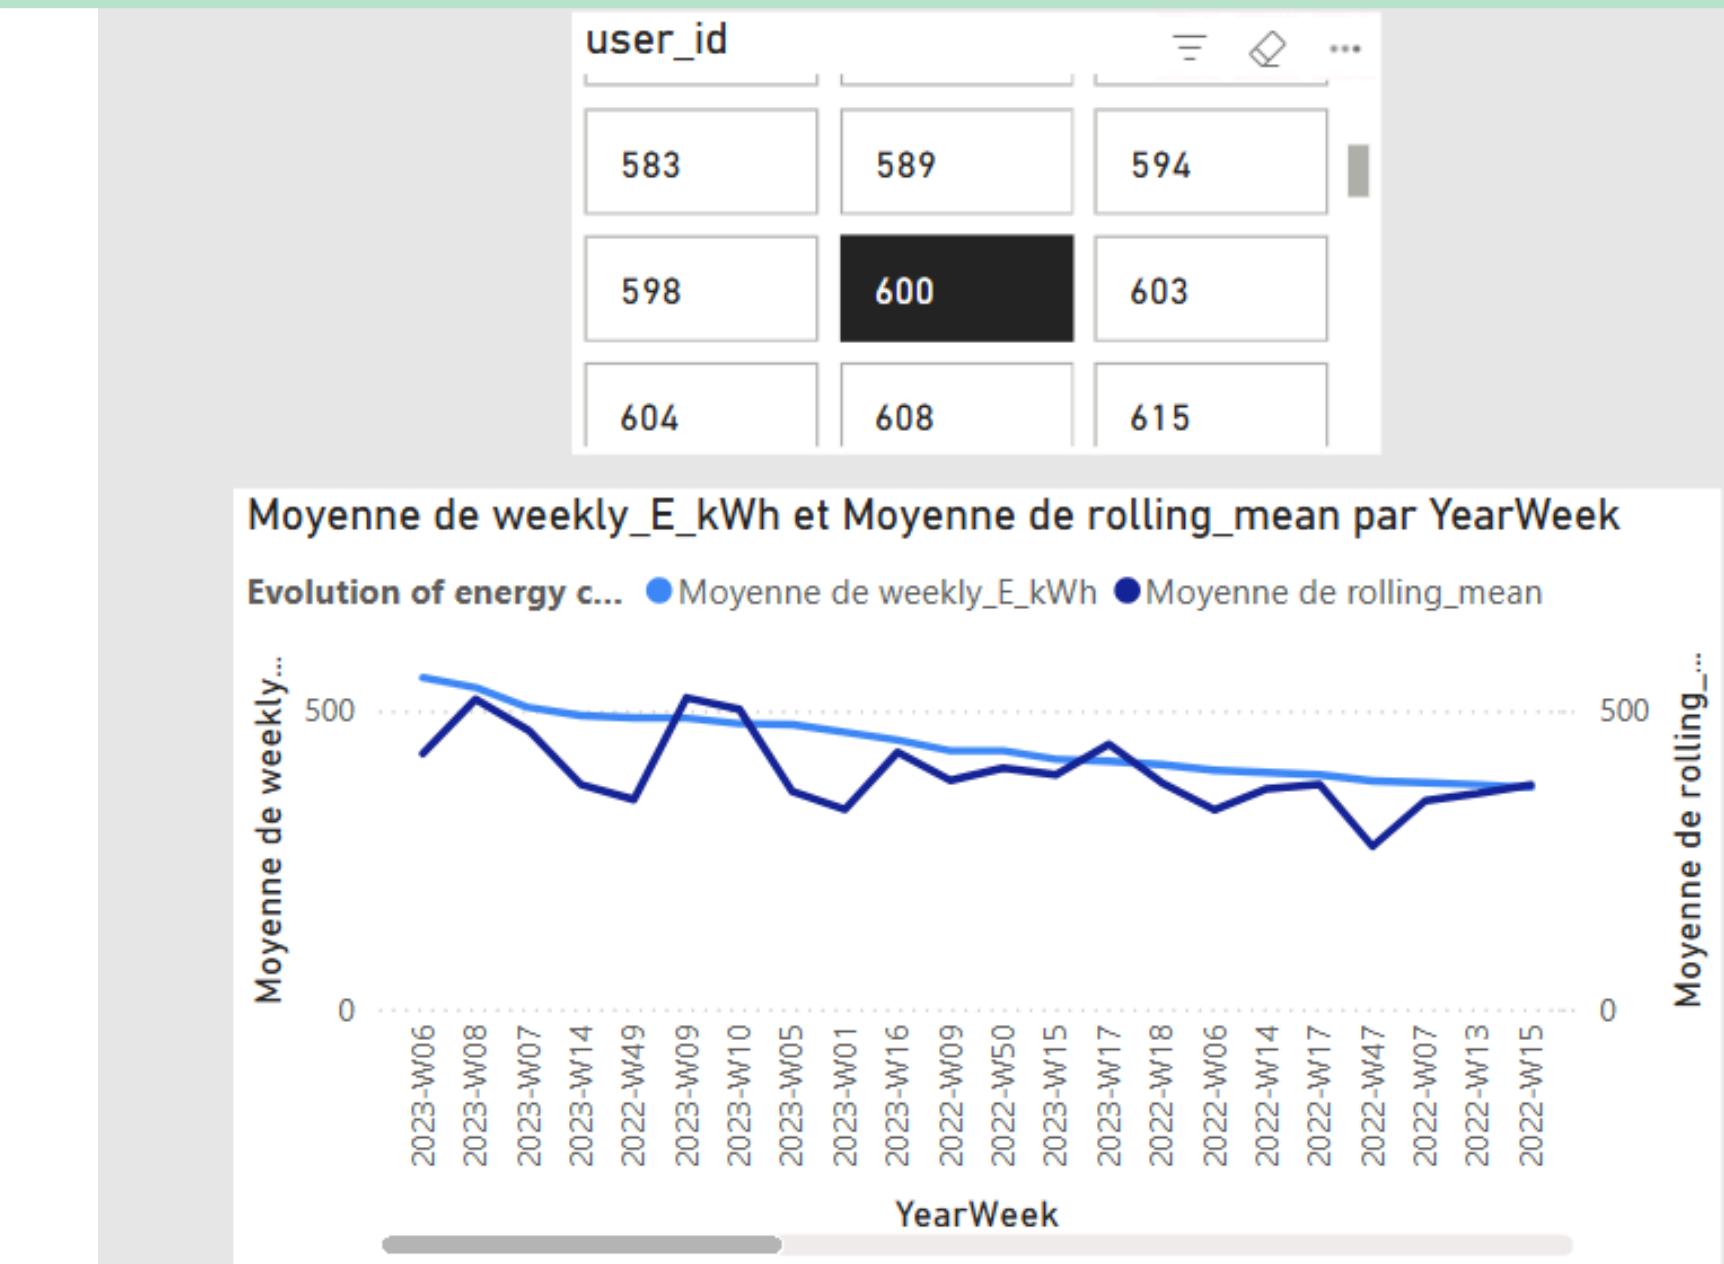
1. Estimated energy consumption

This graph shows how much electricity a user consumes each week, and compares it to a 4-week moving average.

This graph includes an interactive user filter (slicer), allowing us to explore the energy usage of each user individually over time.

- See trends (increasing or decreasing consumption)
- Detect seasonal patterns
- Spot sudden changes in usage

Energy Consumption Over Time



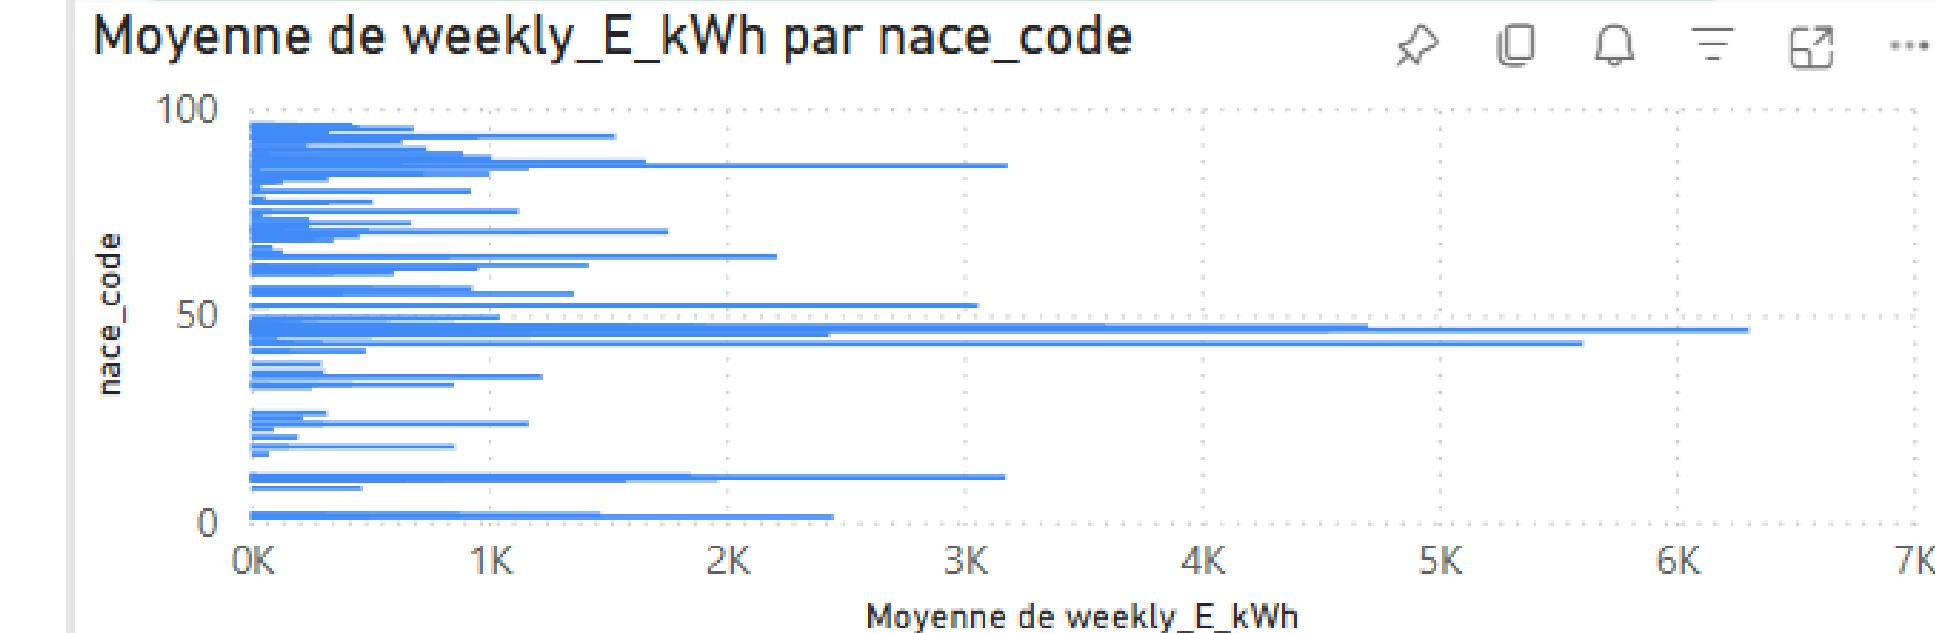
PHASE 3: CREATING A DASHBOARD

1. Estimated energy consumption

It shows the average weekly energy consumption per nace_code, which corresponds to the type of economic activity (industry, education, services, etc.).

- Compare energy usage by sector
- Identify which activity types consume the most electricity
- Help prioritize energy-saving strategies by industry

Energy Use by Activity Type



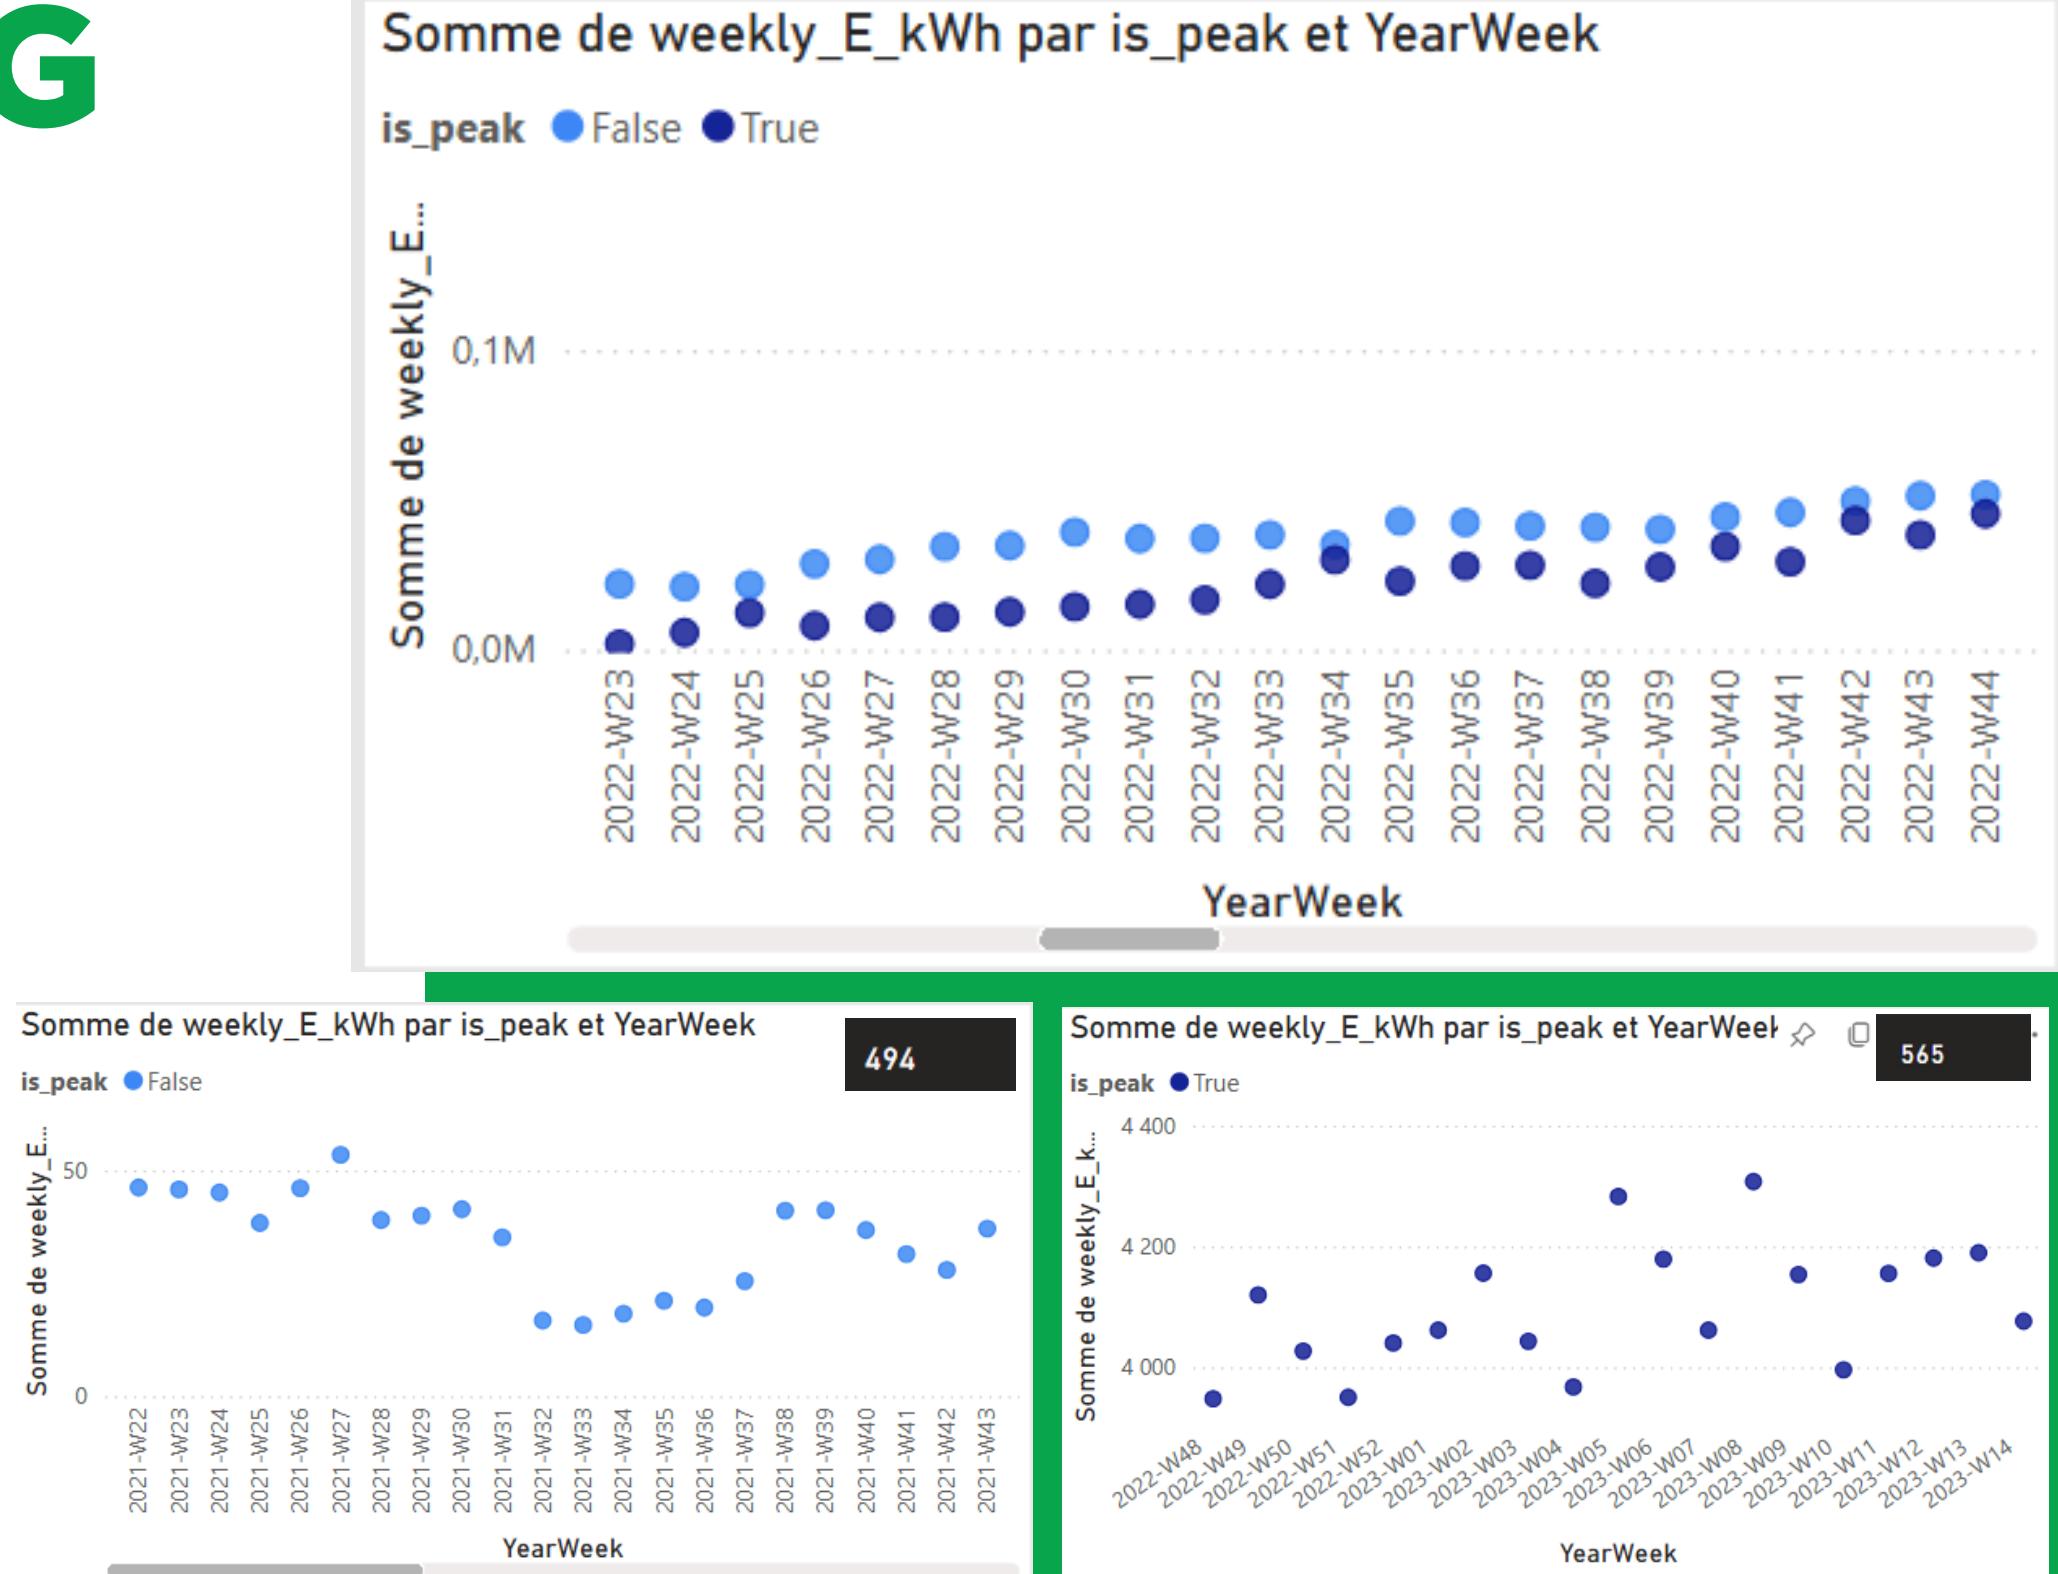
PHASE 3: CREATING A DASHBOARD

2. Detection of consumption peaks

This graph displays only the weeks where energy peaks happened, as dots on a chart.

- Find when high consumption occurred
- Spot outliers and unusual behavior
- Understand which users or weeks are critical

Peak Detection



PHASE 3: CREATING A DASHBOARD

2. Detection of consumption peaks

This table lists all the weeks with a peak, along with key information:
user_id, site_id, YearWeek, nace_code, etc.

-
- Analyze peak cases in detail
 - Identify sectors or buildings using too much energy
 - Easily filter by user, year, or activity

Dynamic Table of Peaks

Dynamic peak table (peak is true)

user_id	YearWeek	weekly_E_kWh	site_id	nace_code	department_encoded	...
501	2023-W23	4 488,75	536	55,10	3	
501	2023-W24	4 242,18	536	55,10	3	
501	2023-W25	3 758,12	536	55,10	3	
565	2022-W48	3 948,00	3843	47,30	5	
565	2022-W49	4 119,87	3843	47,30	5	
565	2022-W50	4 027,04	3843	47,30	5	

PHASE 3: CREATING A DASHBOARD

3. Classification or clustering results

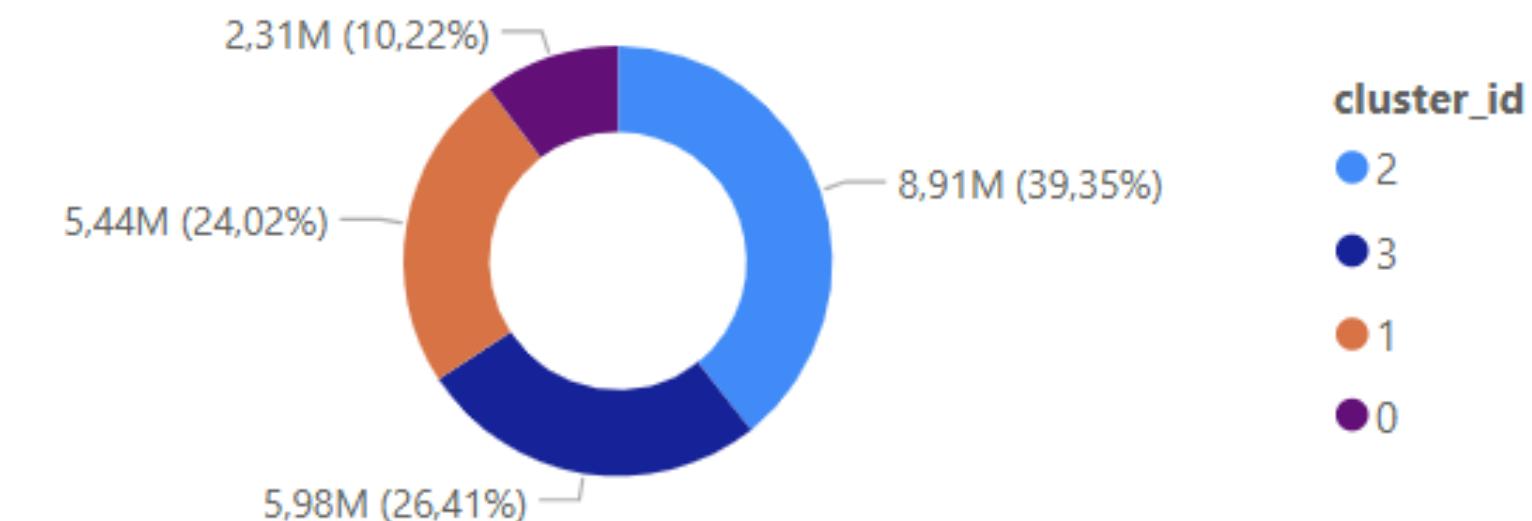
This donut chart displays the total energy consumption per cluster, based on KMeans clustering.

Each cluster_id groups users with similar energy usage behavior, influenced by indoor conditions such as temperature and humidity.

- Understand how consumption is distributed across user groups
- See which cluster is the most energy-intensive
- Use it to target actions by behavioral profiles

Energy Consumption by Consumption Profile (Cluster)

Somme de weekly_E_kWh par cluster_id

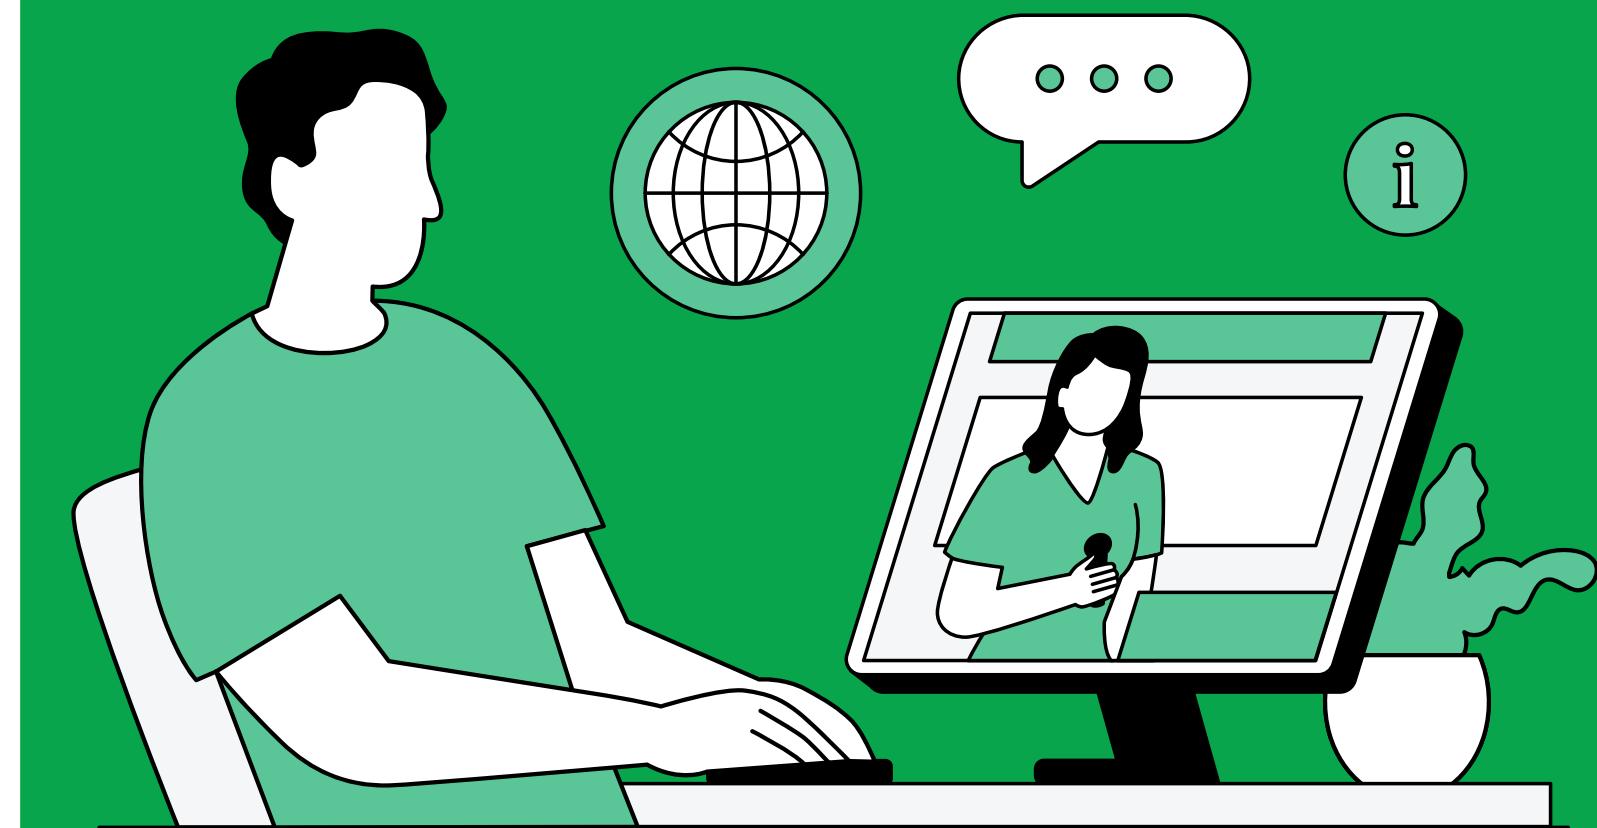


PHASE 3: CREATING A DASHBOARD

In conclusion, this dashboard allowed us to explore and better understand energy consumption patterns through key visualizations.

We analyzed weekly consumption trends, detected consumption peaks, and identified user behavior clusters.

These insights can support targeted energy-saving strategies and better decision-making across sectors.



CONCLUSION: RECOMMENDATIONS

- **Target High-Consumption Activities :**
Focusing on energy audits and efficiency improvements in high-demanding sectors
- **Regional Prioritization :**
Targeted regional energy-saving programs, such as renewable energy incentives, or awareness campaigns in the most consuming areas(like Reunion Island)
- **Predictive Energy Monitoring :**
integrate the Random Forest model into an energy monitoring system to forecast consumption in real time and identify unusual usage before it becomes costly.
- **Automated Anomaly Detection :**
Using Isolation Forest and K-Means models to flag anomalies, such as sudden spikes or abnormal temperature readings
- **Encourage Energy Awareness :**
Share model insights with operational teams to raise awareness



OUR TEAM



ISMAEL



LISA



NGOULAYE



LEINA



KYLIE



ILAN



DJIBRIL



**THANK YOU
FOR YOUR ATTENTION !**